

THE AMD DEPOSITION

A PhD Due Diligence Study of the World's Most Credible NVIDIA Challenger

CPU Dominance, AI GPU Challenge, Xilinx Acquisition & Business Model Validity

August 2026 | Blue Lotus Research | PhD Due Diligence Series

RESEARCH METHODS: *Qualitative (business model, strategy, management) · Quantitative (financial analysis, market data, revenue modelling) · Ethnographic (customer behaviour, developer ecosystem, enterprise procurement) · Superforecasting (10 probability-weighted scenarios) · Adversarial (bull/bear balance sheet per business dimension)*
EPISTEMIC STATUS: *All forward-looking statements labeled as analytical judgment. Financial data sourced from AMD 10-K filings (EDGAR), earnings releases, and market intelligence. Market share data from IDC, Jon Peddie Research, and Mercury Research. No insider information; all public sources cited.*

PART I — THE AMD BUSINESS MODEL: FROM NEAR-DEATH TO CHALLENGER-IN-CHIEF

Advanced Micro Devices (AMD) is the most remarkable turnaround story in semiconductor history. In 2015, AMD's stock traded at \$1.61. The company had \$2.7B in debt, negative book value, was losing market share in every product category, and many analysts considered it a terminal case. In August 2014, AMD's board made a decision that would change everything: it hired Lisa Su as Chief Executive Officer. What followed over the next decade was a systematic strategic transformation — the Zen CPU architecture, the chiplet manufacturing innovation, the Radeon and Instinct GPU lines, and the \$49B Xilinx acquisition — that rebuilt AMD from the ground up. By August 2026, AMD generates \$25B+ in annual revenue, is the leading x86 server CPU challenger to Intel, earns \$12B+ from its data center segment alone, and is the only credible alternative to NVIDIA for AI accelerators at enterprise scale. The question is no longer whether AMD survived — it is how durable its moat is.

1.1 The Fabless Model — Design Without Manufacturing

AMD is a fabless semiconductor company: it designs chips but outsources all fabrication to third-party foundries, primarily TSMC. This model differs fundamentally from Intel's integrated device manufacturer (IDM) model — Intel designs and manufactures its own chips — and creates both advantages and dependencies for AMD. AMD's fabless structure means it carries none of TSMC's \$60B+ annual CapEx burden, allowing it to deploy capital to R&D and design rather than wafer fabs. But it creates complete manufacturing dependency on TSMC — the same Taiwan geopolitical concentration risk shared with NVIDIA.

- ▶ AMD uses TSMC N4/N5 for EPYC 4th gen Genoa CPUs and Instinct MI300X AI accelerators
- ▶ AMD uses TSMC N3 for EPYC 5th gen Turin CPUs (Zen 5 architecture, 2024)
- ▶ AMD uses GlobalFoundries for I/O dies (mature node, lower cost) in chiplet designs
- ▶ Zero manufacturing CapEx: AMD's annual CapEx is ~\$200–300M vs TSMC's \$60–64B
- ▶ Fabless ROIC advantage: AMD returns on invested capital directly from IP/design, not fab infrastructure
- ▶ Manufacturing risk: 100% of AMD's advanced chips depend on TSMC Taiwan — identical geopolitical exposure to NVIDIA

1.2 The Chiplet Architecture Innovation — AMD's Manufacturing Masterstroke

AMD's most consequential engineering decision was the chiplet disaggregation strategy pioneered with EPYC Rome (2019) and perfected with EPYC Milan, Genoa, and Turin. Instead of manufacturing a monolithic processor die (which scales poorly and increases defect probability exponentially with area), AMD disaggregated its chips into: (1) small compute chiplets (CCDs — Core Complex Dies) manufactured at the most advanced node available (N5, N4, N3), and (2) a larger I/O die manufactured at a mature, lower-cost node (GlobalFoundries 12nm). This allows AMD to use TSMC's most expensive leading-edge capacity for the logic that benefits most from it

(compute cores) while using cheaper mature nodes for the I/O infrastructure that does not. The result: AMD's chiplet architecture delivers higher performance-per-watt, higher yields, and lower cost than Intel's monolithic approach — and it is the underlying reason AMD's EPYC has outperformed Intel Xeon for 5+ years.

1.3 The Four Business Segments

Segment	FY2024 Revenue	YoY Growth	Key Products	Primary Market	Margin Profile
Data Center	\$12.565B	+94% YoY	EPYC server CPUs (Genoa, Turin); Instinct MI300X/MI350X AI accelerators	Cloud hyperscalers (AWS, Azure, Google); enterprise data centres	Highest margin — premium ASPs at leading edge
Client	\$6.887B	+52% YoY	Ryzen 7000/9000 desktop + laptop CPUs; Ryzen AI; Ryzen AI Max+	OEM (HP, Dell, Lenovo, ASUS); DIY retail; AI PC segment	Good margins; consumer price sensitivity limits premium
Gaming	\$2.607B	-58% YoY	Radeon RX 7000/9000 consumer GPUs; PlayStation 5 + Xbox APUs (semi-custom)	Consumer gaming; Sony/Microsoft console supply contracts	Lower margins; PS5/Xbox contracts declining as consoles age; GPU share loss to NVIDIA
Embedded	\$3.726B	+8% YoY (recovering from trough)	Xilinx Versal AI Core; Zynq UltraScale+; Kria; Virtex/Artix FPGAs	Telecom (5G infrastructure); industrial automation; automotive; aerospace	Historically high margins; recovering from 2022–2024 inventory correction

SOURCE: AMD FY2024 Annual Report (EDGAR) / AMD Q4 FY2024 earnings press release (Jan 2025) / AMD investor relations segment breakdown

1.4 The Lisa Su Transformation — Management Quality as Competitive Advantage

CEO Lisa Su is widely regarded as one of the most capable semiconductor executives in the industry — a distinction earned through a decade of execution, not reputation. Su joined AMD in 2012, became CEO in 2014, and immediately dismantled the AMD of the previous decade: she killed the APU-only strategy, refocused on high-performance computing, greenlit the Zen architecture (designed by Jim Keller), and restructured the balance sheet. The strategic coherence — every decision from Zen architecture to chiplet manufacturing to Xilinx acquisition to MI300X AI accelerator fits a single narrative — distinguishes AMD's management quality from Intel's decade of strategic drift. AMD's market capitalisation went from ~\$2B (2015) to ~\$200–250B (2026) under Su's tenure, delivering approximately 100–125× shareholder return. Management quality is not a moat by itself, but at AMD it is an underrated competitive advantage.

"The strategy was very simple: we were going to go back to basics. We were going to focus on the best possible products for the high-performance market."

— Lisa Su, AMD CEO — on the post-2014 strategic pivot

FINDING I-A: AMD IS A GENUINE STRUCTURAL CHALLENGER — NOT A NICHE ALTERNATIVE

AMD's business model has evolved from a permanent second-place challenger to a genuine platform company with structural advantages in CPUs (chiplet architecture, TSMC node parity) and a credible,

if software-constrained, AI GPU business. The fabless model gives AMD a CapEx efficiency that Intel cannot match; the Zen architecture executes reliably on a two-year cadence; the MI300X has been purchased by Microsoft Azure and Oracle at scale — proving the AI GPU thesis is not merely promotional. The open questions are not whether AMD is viable — it clearly is — but whether the ROCm software ecosystem can achieve CUDA parity, whether the Xilinx/embedded segment delivers its strategic promise, and whether the TSMC/Taiwan supply chain concentration creates shared risk with NVIDIA.

Evidence: AMD FY2024 revenue: \$25.785B (+13.7%) · Data center: \$12.565B (+94%) — one year of MI300X impact · EPYC market share: ~25–30% server CPU (Mercury Research 2025) · AMD stock: \$1.61 (2015 low) → ~\$120–160 range (2026) · Management tenure: Lisa Su CEO since 2014 — longest-tenured major semiconductor CEO

PART II — FINANCIAL PERFORMANCE: THE DATA CENTER TRANSFORMATION

AMD's financial profile has undergone a structural shift since 2022. The company that once earned the majority of its revenue from consumer CPUs and gaming now generates nearly half its revenue from the data center segment alone. The FY2024 data center figure — \$12.565B, up 94% from FY2023's \$6.5B — is primarily attributable to one product: the Instinct MI300X AI accelerator. This concentration of one new product driving 94% growth in a segment is both AMD's greatest opportunity and its most important monitoring point: if MI300X demand disappoints or NVIDIA widens its software moat, the data center trajectory reverses sharply. The client segment recovery (+52%) confirms Ryzen remains competitive. The gaming decline (-58%) reflects both NVIDIA's consumer GPU dominance and the natural end-of-cycle decline in PS5/Xbox revenue as those console generations age toward replacement.

2.1 Revenue Waterfall — FY2022 to FY2024

Metric	FY2022	FY2023	FY2024	Q1 2025	FY2025 (est.)	FY2026 (est.)
Total Revenue	\$23.601B	\$22.680B (-3.9%)	\$25.785B (+13.7%)	\$7.438B (+36% YoY)	~\$29–32B	~\$34–42B
Data Center	\$6.043B (est.)	\$6.496B	\$12.565B (+94%)	\$3.700B	~\$18–22B	~\$22–30B
Client	\$6.188B	\$4.530B (-27%)	\$6.887B (+52%)	~\$1.9B	~\$7–8B	~\$8–9B
Gaming	\$6.805B	\$6.166B (-9%)	\$2.607B (-58%)	~\$0.7B	~\$2.5–3.0B	~\$2–3B
Embedded (Xilinx)	\$4.565B	\$5.449B → trough \$3.4B	\$3.726B (+9% recovery)	~\$0.8B	~\$3.5–4.5B	~\$4–6B
Gross Margin (GAAP)	50.5%	46.1%	52.7%	~50%	~52–54%	~53–56%
Net Income (GAAP)	\$1.264B	-\$0.177B (loss)	\$1.641B	~\$0.5B	~\$2.5–4B	~\$4–7B
Net Income (Non-GAAP)	\$4.583B	\$4.050B	\$5.433B	~\$1.57B	~\$7–9B	~\$9–14B

SOURCE: AMD FY2024 Annual Report (EDGAR) · AMD Q4 FY2024 + Q1 FY2025 earnings releases · Wall Street consensus estimates (Goldman Sachs, JPMorgan, Morgan Stanley, Blue Lotus base) for FY2025-FY2026

2.2 The GAAP vs Non-GAAP Gap — Understanding the Real Economics

AMD's GAAP vs Non-GAAP earnings divergence is significant and requires explicit analysis. FY2024 GAAP net income: \$1.641B. Non-GAAP net income: \$5.433B. The \$3.8B gap consists primarily of: (1) stock-based compensation (~\$1.4B) — real economic cost, dilutes shareholders; (2) amortization of intangible assets (~\$2.0B) — Xilinx acquisition goodwill amortization; (3) restructuring and one-time items (~\$0.4B). The Xilinx amortization is a genuine accounting artifact

of the \$49B acquisition price — it depresses GAAP earnings significantly but does not represent a cash outflow. Stock-based compensation (SBC) is real dilution but is standard for semiconductor companies competing for engineering talent. Blue Lotus uses non-GAAP net income for operational performance comparison while flagging SBC as a genuine shareholder cost.

AMD FY2024 Financial Bridge (GAAP → Non-GAAP):		
GAAP Net Income:		\$1.641B
+ Stock-Based Compensation (SBC):	+\$1.392B (dilutive — 5,400+ engineers/executive)	
+ Xilinx intangible amortization:	+\$1.973B (accounting; not cash)	
+ Restructuring + other adjustments:	+\$0.427B	
= Non-GAAP Net Income:		\$5.433B
Non-GAAP Gross Margin FY2024: 53.3% (vs GAAP 52.7%)		
Non-GAAP EPS FY2024: \$3.31 per diluted share		
AMD Market Cap (Aug 2026 est.): ~\$200–250B → P/E (Non-GAAP FY2024): ~37–46x		
Note: P/E on FY2025/FY2026 estimates is lower as earnings accelerate		
All AMD financial figures from EDGAR 10-K FY2024. Non-GAAP adjustments are AMD's disclosed reconciliation — Blue Lotus presents both, with non-GAAP used for comparability to peers.		

2.3 Balance Sheet and Capital Allocation

- Cash and equivalents: ~\$4.3B (Q4 FY2024); no significant near-term debt maturities
- Total debt: ~\$1.7B long-term; net cash positive — strong balance sheet post-Xilinx integration
- Annual CapEx: ~\$200–300M (fabless model; TSMC handles manufacturing CapEx)
- Free cash flow (GAAP): \$1.6B FY2024; \$5.4B non-GAAP adjusted FCF
- Share buyback: \$3.4B repurchased FY2024; \$6B remaining authorisation
- No dividend: AMD reinvests FCF into R&D (\$5.8B in FY2024 = ~22% of revenue — highest in AMD history)
- Goodwill: ~\$24.2B (predominantly Xilinx \$49B acquisition value); no impairment as of FY2024

2.4 Revenue Quality Analysis — How Durable Is 94% Data Center Growth?

The 94% YoY growth in AMD's data center segment in FY2024 is extraordinary — but its durability depends almost entirely on MI300X/MI350X AI accelerator adoption trajectory. AMD management guided for data center revenue growth to continue in 2025, with MI300X revenue embedded in the data center segment growing toward \$10B+ annually. The critical analytical question is whether AI GPU revenue is sustainable or reflects a one-time adoption surge. The evidence suggests sustained demand — multiple hyperscalers (Azure, Oracle) have committed to multi-year MI300X/MI350X infrastructure build-outs — but the growth rate will normalise toward 20–35% as the installed base grows. EPYC server CPU revenue is the more durable, lower-growth component of data center: at ~25–30% server market share and growing, EPYC contributes a reliable \$6–8B annually with ~5–8% secular growth from cloud infrastructure expansion.

[EDGAR] Query: "AMD revenue data center segment growth FY2024..." Score:49.7

TICKER: AMD | COMPANY: ADVANCED MICRO DEVICES INC | FORM: 10-Q | DATE: 2025-08-06 | SECTION: Item 2 — MANAGEMENT'S DISCUSSION AND ANALYSIS OF FINANCIAL CONDITION and Gaming segment revenue primarily driven by strong demand for the latest "Zen 5" AMD Ryzen™ processors and higher semi-custom sales, and an increase in Data Center segment revenue primarily driven by growth in AMD EPYC™ CPU sales. Embedded segment revenue decreased as end market

demand remained mixed. Gross margi...

FINDING II-A: AMD'S FINANCIAL TRANSFORMATION IS STRUCTURALLY REAL — BUT MI300X TRAJECTORY IS THE KEY VARIABLE

AMD's FY2024 financial results demonstrate a structurally transformed company. Data center revenue at \$12.565B (+94%) is primarily driven by MI300X AI accelerator adoption — a single product generating multi-billion dollar revenue within two years of launch. The EPYC CPU business provides a durable \$6–8B base with predictable growth. The gross margin expansion (52.7% FY2024, vs 46.1% FY2023) reflects product mix shift toward higher-ASP data center silicon — the same dynamic TSMC experiences as it ramps N2. The risks are concentrated: if MI300X/MI350X adoption slows (NVIDIA software moat wins), data center revenue growth collapses. If EPYC loses market share (Intel 18A executes), the stable cash engine underperforms. Either scenario alone is manageable; both together would be a serious problem. Currently, neither is imminent.

Evidence: FY2024 data center: \$12.565B (+94% YoY) · FY2024 non-GAAP GM: 53.3% · FY2024 R&D: \$5.8B (~22% of revenue) · Net cash positive: ~\$4.3B cash vs \$1.7B debt · MI300X customers: Microsoft Azure 5,000+ units (est.), Oracle AI, Meta training clusters

PART III — THE CPU BATTLE: HOW AMD WENT FROM 0% TO 25–30% SERVER MARKET SHARE IN EIGHT YEARS

The AMD CPU story is the most improbable corporate turnaround in semiconductor history. In 2016, AMD's server CPU market share was effectively zero — below measurement threshold. Intel's Xeon owned the data centre, and AMD's Bulldozer architecture was so power-inefficient and so computationally weak that major cloud providers had quietly removed AMD options from their infrastructure planning roadmaps entirely. Eight years later, AMD's EPYC processors hold approximately 25–30% of the server CPU market by revenue — a share that continues to expand — and Intel has been forced into a defensive posture on its own home turf. The transformation was not marketing. It was architecture: the chiplet disaggregation strategy that Intel dismissed as inferior, and that proved to be the manufacturing insight of the decade.

3.1 The Zen Architecture Revolution — From Bulldozer to Zen 5

The Bulldozer architecture (2011–2015) was AMD's near-death experience in computing. Its shared floating-point units between paired integer cores produced a design that consumed excessive power, ran hot, and underperformed Intel's Sandy Bridge by 10–30% in most enterprise workloads. AMD's server share collapsed from approximately 26% in 2006 to sub-1% by 2016. Lisa Su, appointed CEO in October 2014, made a binary bet: kill Bulldozer, invest the remaining R&D budget in a clean-sheet microarchitecture designed around modern IPC (instructions per clock) efficiency. The result was Zen. Zen 1 (2017) delivered a 52% IPC improvement over Bulldozer in a single generation — one of the largest single-generation IPC leaps in x86 history. Each subsequent Zen generation — Zen 2 (2019), Zen 3 (2021), Zen 4 (2022), Zen 5 (2024) — has maintained competitive or superior IPC vs Intel's equivalent generation.

Generation	Launch	Process Node	IPC vs Prior Gen	Key EPYC Product	Peak Cores (Server)	Server Share Impact
Bulldozer/Excavator	2011–2015	32nm–28nm	Baseline (negative — underperformed)	Opteron 6300 series	16 cores	Collapsed from 26% → <1%
Zen 1	2017	GlobalFoundries 14nm	+52% IPC vs Bulldozer	EPYC Naples (7001)	32 cores	Initial re-entry; ~5% server share
Zen 2	2019	TSMC 7nm + GF 14nm I/O	+15% IPC; chiplet-interposer debut	EPYC Rome (7002)	64 cores	~7–8% server share; first major cloud wins
Zen 3	2021	TSMC 7nm CCD + GF 12nm I/O	+19% IPC; unified L3 per CCD	EPYC Milan (7003)	64 cores + 3D V-Cache	~10–15% server share; AWS, Google wins
Zen 4	2022–	TSMC 5nm CCD	+13% IPC; AVX-	EPYC	128 cores	~18–22%

	2023	+ GF 6nm I/O	512 native; DDR5	Genoa (9004); Bergamo (9754)	(Bergamo)	server share; density leadership
Zen 5	2024– 2025	TSMC N3 CCD + GF 6nm I/O	+16% IPC; wider AVX-512; ML inference	EPYC Turin (9965X)	192 cores	~25–30% server share; competitive with Xeon on all metrics

SOURCE: AMD Investor Presentations 2017–2025 · AnandTech Zen architecture analysis · Mercury Research server CPU share · Serve The Home EPYC benchmark database · TSMC process node assignments (EPYC CCD roadmap)

3.2 The Chiplet Innovation — AMD's Manufacturing Moat

The chiplet disaggregation strategy is AMD's defining architectural innovation — and the one that Intel most spectacularly underestimated. Traditional monolithic processors place all CPU functions (compute cores, memory controller, I/O interfaces, PCIe lanes) on a single die. This approach has a critical manufacturing economics problem: as die area grows, yield falls exponentially — a 400mm² die at a given defect density yields far fewer functional parts than two 200mm² dies combined. AMD solved this by disaggregating the die. Compute Core Chiplets (CCDs) — the small, high-value dies that contain the Zen CPU cores and L2/L3 cache — are fabbed at TSMC's leading node (N3 for Zen 5). The I/O Die — which contains memory controllers, PCIe controllers, and interconnect logic — is fabbed at GlobalFoundries 6nm or 12nm, which is cheaper and does not need the leading node. This split gives AMD the leading-node performance for compute without paying leading-node economics for I/O infrastructure — a structural cost advantage Intel's monolithic Xeon cannot match without adopting the same architecture (which Intel is now doing with its own tile-based Xeon designs, explicitly acknowledging AMD's approach was correct).

AMD Chiplet Economics (EPYC Turin — Zen 5 192-core):

Compute Core Chiplets (CCDs): 12× Zen 5 CCD dies at TSMC N3
Each CCD: 16 Zen 5 cores + 32MB L3 cache; small die area → high N3 yield
12 CCDs × 16 cores = 192 cores total

I/O Die: 1× IOD at GlobalFoundries 6nm
Contains: 12× LPDDR5/DDR5 memory channels; PCIe 5.0 lanes; Infinity Fabric
GF 6nm: ~\$X/wafer vs TSMC N3: ~\$Y/wafer (N3 >3× more expensive)

Result: AMD pays TSMC N3 prices only for the cores (where it matters);
pays GF 6nm prices for I/O (where node doesn't matter).
Intel pays Intel 4 (or TSMC N3/N18A) prices for the ENTIRE monolithic die.

Manufacturing Advantage: AMD's chiplet approach gives it superior gross margin

structure for equivalent performance vs Intel's monolithic approach.

I/O Die wafer costs are not publicly disclosed; relative cost advantage estimated from die area analysis and node pricing reports
(SemiAnalysis, Doug O'Laughlin)

- ▶ 3D V-Cache: AMD stacks additional L3 cache directly on top of CCD dies via TSMC SoIC — adds 64–192MB L3 per CCD; critical for gaming + HPC + AI inference workloads
- ▶ EPYC Bergamo: 128-core cloud-optimised variant (Zen 4c dense cores); targets hyperscaler density requirements; directly competes with ARM (Graviton/Ampere) on core-count economics
- ▶ Multi-chip Module (MCM) interconnect: AMD Infinity Fabric connects CCDs at 32–64 GB/s bandwidth; coherent shared memory across all 192 cores in Turin
- ▶ Packaging: AMD uses TSMC CoWoS and TSMC SoIC for advanced stacking; same supplier as NVIDIA — creates shared allocation risk

3.3 The Intel Xeon Battle — Market Share Data and Competitive Positioning

AMD's EPYC versus Intel Xeon is the defining competition in enterprise computing. Intel held a structural monopoly in server CPUs from 2006 to 2017 — not because of superior products, but because AMD had no credible alternative. EPYC Naples (Zen 1, 2017) was AMD's breach: the first credible enterprise alternative. Each subsequent EPYC generation has gained share. By Q1 2026, AMD EPYC holds approximately 25–30% of the server CPU market by revenue — with some cloud providers (Google, Microsoft Azure, Oracle) deploying EPYC at rates that suggest 30–40%+ of their new CPU server deployments. Intel's Xeon response has been defensive: the Sapphire Rapids (4th Gen Xeon, 2023) launch was delayed multiple times; Granite Rapids (5th Gen Xeon, 2024) improved but did not close AMD's core-count or memory-bandwidth lead. Intel's next bet — Clearwater Forest (Intel 18A process, 2025–2026) — represents an attempt to out-chiplet AMD using Intel's own process, but execution risk is high given Intel Foundry's yield challenges.

Metric	AMD EPYC Turin (Zen 5, 2024)	Intel Xeon Granite Rapids (2024)	Winner
Max Core Count	192 cores (EPYC 9965X)	128 cores (Granite Rapids)	AMD (+50%)
Process Node	TSMC N3 (CCD) + GF 6nm (I/O)	Intel 3 (Intel Foundry)	TIE — Intel 3 vs N3 comparable
Memory Channels	12× DDR5-6000	8× DDR5-5600	AMD (+50% channels; +43% BW)
PCIe Lanes	160× PCIe 5.0	96× PCIe 5.0	AMD (+67% lanes)
L3 Cache (max)	384MB (12 CCDs × 32MB) + 3D V-Cache option	320MB (Granite Rapids)	AMD (base); AMD++ (with V-Cache)
TDP Range	120W–500W	60W–500W	TIE — comparable envelope
Single-Thread Perf (SPECint est.)	~Competitive / slight AMD lead	~Competitive / closing gap	TIE to AMD slight
Virtualisation Density	192 vCPUs (superior for cloud)	128 vCPUs	AMD (+50% VM density)
List Price (flagship)	~\$14,000 (EPYC 9965X 192c)	~\$17,000 (Xeon Platinum 8592+)	AMD (~18% price; superior value)
Memory Bandwidth (TB/s)	~0.9 TB/s (12-channel × 76.8 GB/s)	~0.61 TB/s (8-channel)	AMD (+48% bandwidth)
AI/ML Inference Built-in	Zen 5 AVX-512 VNNI; BFLOAT16	AMX matrix extensions	Intel (AMX advantage for specific workloads)
Server Market Share 2024	~25–30% revenue	~70–75% revenue	Intel (holds volume; AMD gaining)

SOURCE: AnandTech/Chips and Cheese EPYC Turin technical analysis 2024 · Intel Granite Rapids launch specs · Mercury Research CPU market share Q4 2024 · Serve The Home EPYC Genoa/Turin benchmarks · AMD Q4 2024 earnings call (Dr. Lisa Su commentary on server CPU trajectory)

FINDING III-A: AMD'S CPU ARCHITECTURE ADVANTAGE IS NOW STRUCTURAL, NOT CYCLICAL

AMD's chiplet disaggregation is no longer a temporary advantage — it is a permanent architectural lead that Intel is scrambling to replicate. The core count leadership (192 vs 128), memory bandwidth

advantage (12-channel vs 8-channel), and price-performance superiority are not aberrations of a single generation — they reflect an architectural philosophy that scales better than Intel's monolithic approach. Intel's response (Clearwater Forest, Sierra Forest, Granite Rapids DC) all adopt Intel-version chiplets, acknowledging AMD was correct. AMD's server share trajectory — from 0% in 2016 to 25–30% in 2025, expected 30–35% by 2027 — represents a structural rebalancing of the x86 server ecosystem. The question is not whether AMD takes more share — it is how much Intel can defend.

Evidence: AMD EPYC server revenue FY2024: ~\$10B+ of the \$12.565B Data Center segment (est.) · Mercury Research Q4 2024: AMD ~25–30% server CPU market share · EPYC adoption: Google, Microsoft Azure, Oracle, Meta all cite EPYC as 'preferred CPU platform' · Intel Xeon revenue: fell 15–20% YoY in 2023–2024 as AMD captured incremental cloud demand

PART IV — THE AI GPU CHALLENGE: THE BATTLE TO BREAK NVIDIA'S CUDA MONOPOLY

AMD's AI GPU business is the most watched competitive narrative in the semiconductor industry. NVIDIA holds approximately 80–87% of the data centre AI GPU market by revenue — a dominance built not on hardware alone but on the CUDA software ecosystem: 4,000+ optimised applications, more than 1 million developers, and 15 years of PyTorch-to-CUDA workflow lock-in that makes switching prohibitively expensive for most enterprises. AMD's challenge is therefore not primarily a silicon challenge — the MI300X hardware is genuinely competitive with the H100 on several dimensions — but a software ecosystem challenge. Every day that PyTorch's default path is CUDA, every day that a machine-learning researcher writes CUDA-optimised kernels, is another day AMD's ROCm ecosystem falls further behind. This is the central strategic question for AMD's AI GPU business: can the hardware be good enough, long enough, to force a software ecosystem inflection that breaks CUDA's gravitational pull?

4.1 The MI300X — AMD's Best AI Accelerator Argument

The Instinct MI300X, launched in late 2023 and ramping through 2024, is AMD's strongest argument that NVIDIA's hardware monopoly is not insurmountable. The MI300X is an Accelerated Processing Unit (APU) — a novel architecture that places 8 CDNA3 GPU dies and 4 CPU-class dies on a single package using TSMC's advanced 3D packaging, all sharing a unified 192GB HBM3 memory pool. This architecture delivers a key inference advantage: the ability to fit an entire 70B-parameter LLaMA-class model in a single GPU's memory, whereas NVIDIA's H100 (80GB HBM3) requires multi-GPU coordination for the same model. For inference workloads — where latency matters more than raw throughput — this memory advantage translates to real commercial value.

Specification	AMD MI300X	AMD MI350X (2025)	AMD MI400 (2026 est.)	NVIDIA H100 SXM5	NVIDIA B200
Process Node	TSMC N5P (GPU) + N6 (CPU)	TSMC N3 (est.) / N5P enhanced	TSMC N3 (expected)	TSMC N4 (A100 successor)	TSMC N4P
GPU Architecture	CDNA3	CDNA3.5 / CDNA4	CDNA4 (next-gen)	Hopper	Blackwell
HBM Memory	192GB HBM3 (unified pool)	288GB HBM3E (+50% vs MI300X)	TBD; HBM4 expected	80GB HBM3	192GB HBM3E
Memory BW (TB/s)	5.3 TB/s	~8.0 TB/s (est.)	~12+ TB/s (est.)	3.35 TB/s	8.0 TB/s
FP8 Throughput	~2,600 TFLOPS	~5,000 TFLOPS (est.)	~10,000+ TFLOPS (est.)	~3,958 TFLOPS	~9,000 TFLOPS
FP16 Throughput	1,307 TFLOPS	~2,400 TFLOPS (est.)	~5,000 TFLOPS (est.)	989 TFLOPS	4,500 TFLOPS
Packaging	TSMC CoWoS-L (8 GPU + 4 CPU dies)	CoWoS-L enhanced	CoWoS-L (est.)	TSMC CoWoS-S	CoWoS-L
TDP	750W	~750–800W	TBD	700W	1,000W

		(est.)			
List Price (est.)	\$10,000–\$15,000	\$15,000–\$20,000 (est.)	TBD	\$25,000–\$40,000	\$40,000–\$50,000+
Key Strength	192GB unified memory — fits 70B+ LLM in one GPU	288GB HBM3E — fits 140B+ models	Next-gen; target: match B200	CUDA ecosystem; production-proven	Blackwell architecture; CUDA ecosystem
Key Weakness	ROCm vs CUDA ecosystem gap; FP8 lower than H100	ROCm ecosystem still maturing	Competing with NVIDIA Rubin (2027)	Only 80GB — multi-GPU for large models	Higher cost; still TSMC CoWoS constrained
Primary Market	Inference (memory-bound LLM serving)	Training + Inference	Full AI stack	Training (dominant)	Training + Inference

SOURCE: AMD MI300X product brief · AMD Investor Day 2024 AI GPU roadmap · SemiAnalysis MI300X competitive analysis · AMD Q3 2024 earnings: MI300X 'exceeds \$5B revenue run rate' (Lisa Su) · NVIDIA Blackwell product specifications 2024 · Dylan Patel SemiAnalysis B200 memory bandwidth analysis

4.2 Who Is Buying AMD AI GPUs — And Why

AMD's MI300X customer base is concentrated among hyperscalers and cloud providers who have the engineering resources to work through ROCm's compatibility challenges in exchange for the memory and cost advantages MI300X provides. Unlike NVIDIA's broad enterprise and research market penetration, AMD's AI GPU sales are almost entirely to a small number of sophisticated operators. This concentration both validates MI300X's technical capabilities and exposes AMD's continued dependence on customer willingness to invest in ROCm compatibility work.

- Microsoft Azure: Azure ND MI300X v5 instances launched 2024; used for Azure AI inference services and OpenAI model serving; Microsoft publicly cited MI300X's 192GB HBM as the key procurement driver for inference workloads
- Oracle Cloud: OCI Supercluster MI300X configurations; 8,192-GPU clusters for AI training; Oracle cited price-performance and memory capacity as decision factors vs H100
- Meta (Facebook): Internal deployment of MI300X for inference of Llama family models; Meta is also a major ROCm contributor — direct incentive to mature the ecosystem
- Google (DeepMind / Vertex AI): Limited MI300X deployment; Google primarily uses its own TPUs + NVIDIA but has confirmed MI300X evaluation
- Neocloud providers (CoreWeave, Lambda Labs): Primarily NVIDIA-based; some MI300X deployments where customers specifically request inference price-performance
- NVIDIA's own internal estimate: AMD Instinct took ~12–13% of the AI accelerator market by revenue in 2024 — a number AMD has not confirmed but has not disputed

[EDGAR] Query: "AMD MI300X Instinct AI accelerator Microsoft Azure Orac..." Score:54.5

TICKER: AMD | COMPANY: ADVANCED MICRO DEVICES INC | FORM: 10-K | DATE: 2026-02-04 | SECTION: Item 1 — BUSINESS ur data center AI accelerator products was strong as large hyperscale customers, OEMs and ODMs deployed our AMD Instinct™ MI350X Series GPUs. We advanced our AMD AI GPU roadmap to deliver an annual cadence of leadership for AMD Instinct solutions, beginning with the AMD Instinct MI350 Series GPUs in 2025. Beyond GPUs, we launched the 5th Gen AMD EPYC™ family of serv...

4.3 ROCm vs CUDA — The Software Moat That Defines AMD's AI GPU Ceiling

CUDA (Compute Unified Device Architecture) is NVIDIA's parallel computing platform, launched in 2006. Over 18 years, CUDA has become the de facto language of AI compute — not because it was technically superior to all alternatives, but because NVIDIA made it the default target for PyTorch, TensorFlow, and virtually every ML framework, backed by heavy developer tooling investment. The result: 4,000+ CUDA-optimised applications, 1 million+ developers who know CUDA, and a training corpus of CUDA-optimised code that makes writing new ML research faster in CUDA than in any alternative. AMD's ROCm (Radeon Open Compute) ecosystem is the challenger. ROCm is genuinely open-source (MIT and Apache 2.0 licensed — CUDA is proprietary) and has matured significantly since 2022. PyTorch's official support for ROCm backend (added 2020, substantially improved 2022–2024) means that basic training and inference workloads can run on MI300X with minimal code changes. But 'basic workloads' is not the competitive frontier — it is the specialised CUDA kernels, FlashAttention implementations, Triton GPU kernel compilations, and optimised inference runtimes (TensorRT, vLLM CUDA) that represent where the performance gap lives.

Dimension	NVIDIA CUDA	AMD ROCm	Gap Assessment
Age / Maturity	18 years (2006–2024)	~12 years (2012–2024; serious investment since 2019)	CUDA 6+ year structural lead
Developer Ecosystem	1M+ developers trained in CUDA; default ML research target	~50,000–100,000 ROCm-proficient developers (est.)	CUDA 10–20× larger community
Optimised Applications	4,000+ CUDA-optimised applications; every major ML framework	~500–1,000 ROCm-native applications; PyTorch ROCm official	CUDA 4–8× more applications
PyTorch Support	Default / first-class; all features supported	Official ROCm backend; most features supported (exceptions: some specialised ops)	CUDA slight edge; ROCm gap closing rapidly
FlashAttention (critical for LLMs)	FlashAttention-2/3 CUDA: production-ready; fastest attention impl.	Flash Attention ROCm: available but historically slower; 2024 improvements	CUDA lead, but AMD closing in 2024–2025
Triton GPU Kernels	Triton primary target is CUDA; full support	ROCm Triton backend: functional; not all ops optimised	CUDA clear lead
vLLM Inference Framework	vLLM: CUDA-native; full optimisation; production ready	vLLM ROCm: supported; some performance gaps on specific models	CUDA lead
Proprietary ML Libraries	cuDNN, cuBLAS, TensorRT: NVIDIA-maintained, constantly optimised	MIOpen, rocBLAS, MIGraphX: AMD equivalents; maturing but not cuDNN-competitive on all ops	CUDA meaningful lead
Enterprise Support / Tooling	NVIDIA AI Enterprise: full stack; ISV certifications; SLA support	ROCm Enterprise: improving; fewer ISV certifications	CUDA lead
Open Source Licensing	Proprietary (lock-in tool)	Open source (MIT/Apache) — long-term moat-breaker	ROCm strategic advantage (not yet realised at scale)
HPC Scientific Computing	CUDA dominant in all major HPC codes (GROMACS, NAMD, CP2K)	ROCm HPC support: improving; major HPC centres (ORNL Frontier) use ROCm	ROCm credible in HPC; CUDA dominant in ML
Switching Cost Estimate	—	—	Estimate: 6–18 months of engineering work to port CUDA-optimised

			production system to ROCm at enterprise scale
--	--	--	---

SOURCE: NVIDIA CUDA developer ecosystem statistics · AMD ROCm GitHub ecosystem analysis · MLCommons inference benchmark suite ROCm vs CUDA results 2024 · Semianalysis 'CUDA Moat' analysis · PyTorch ROCm support documentation (pytorch.org) · Oak Ridge National Lab Frontier supercomputer (world's first exascale: AMD EPYC + AMD Instinct MI250X with ROCm)

4.4 AMD's Inference Advantage — The Strategic Wedge

Despite the ROCm software gap, AMD has identified and is exploiting a specific market segment where its hardware advantage is compelling enough to overcome the software friction: large-model inference. As transformer-based AI models grew to 70B, 140B, and 405B parameters, GPU memory became the binding constraint. An NVIDIA H100 (80GB HBM3) cannot fit a Llama 3 405B model in a single GPU — it requires multi-GPU tensor parallelism, adding inter-GPU communication latency that degrades inference throughput. The MI300X, with 192GB HBM3 in a unified pool, can fit 70B models natively and serve them with lower latency. This is not a niche advantage — it directly maps to the most commercially deployed AI models (the 70B class) and the highest-margin AI service (inference/API serving). Microsoft's decision to deploy MI300X for Azure AI inference services was explicitly driven by this memory architecture advantage.

AMD MI300X Inference Economics (70B Parameter Model):

Model Size (70B params, FP16): ~140 GB → fits within MI300X 192GB ✓
NVIDIA H100 (80GB HBM3): 140 GB > 80 GB → requires 2× H100 with NVLink

Single MI300X serving cost:

Hardware: ~\$12,000 (est. cloud-amortised); Power: 750W
Tokens/second (70B FP16, est.): ~400–600 tokens/sec
Cost per 1M tokens: ~\$1.20–\$2.50 (vs \$2.50–\$5.00 for 2× H100 config)

Training workloads: NVIDIA H100/B200 remains preferred
(CUDA ecosystem + NVLink bandwidth > MI300X advantages for training)

Verdict: AMD MI300X is ~30–50% cheaper per token than H100 for 70B inference;
comparable or inferior for 7B/13B models (where H100's CUDA ecosystem wins).

Token cost estimates based on Azure ND MI300X v5 instance pricing and publicly disclosed MI300X benchmark throughput; actual costs vary by workload and utilisation rate

4.5 The MI350X and MI400 Roadmap — Can AMD Stay Competitive?

AMD's AI GPU roadmap demonstrates credible intent to remain competitive, but each generation faces the same competitive dynamic: NVIDIA is also shipping a new generation on a similar cadence. The MI350X (2025), with 288GB HBM3E and improved CDNA architecture, upgrades AMD's memory advantage. The MI400 (2026), expected at TSMC N3, targets NVIDIA's Blackwell B200 generation. The critical question is not whether AMD's hardware is competitive — it has been, on select workloads — but whether ROCm matures fast enough that enterprises can deploy MI400 without prohibitive software integration overhead. AMD's open-source ROCm strategy and its willingness to accept PyTorch contributions may be the structural differentiator: if ROCm becomes genuinely first-class in PyTorch (not just 'supported'), the software switching cost falls dramatically, and AMD's hardware economics become compelling at enterprise scale.

- ▶ MI350X (2025): 288GB HBM3E (+50% over MI300X 192GB); CDNA3.5 or CDNA4 architecture; targets training + inference balance; expected volume H2 2025
- ▶ MI400 (2026): Full next-generation CDNA4; TSMC N3 GPU dies; HBM4 integration target; designed to compete with NVIDIA Blackwell/Rubin era products
- ▶ ROCm 6.x (2024–2025): AMD's most significant ROCm investment cycle; FlashAttention-3 ROCm port; improved vLLM ROCm backend; Triton ROCm compiler improvements
- ▶ PyTorch ROCm parity project: AMD engineers contributing directly to PyTorch to ensure ROCm backend achieves performance parity on core operations by 2026
- ▶ AMD+Microsoft partnership: Azure's public MI300X deployment validates AMD as enterprise-grade; every quarter of Azure MI300X production deployment reduces enterprise software risk perception

[NIKKEI] Query: "AMD ROCm software ecosystem CUDA alternative AI machine..." Score:34.5

to establish a software ecology in China that can rival Nvidia's. The American company released its parallel computing platform CUDA in 2006, including a development toolkit that enabled anyone with a laptop equipped with an Nvidia GPU to develop software. Over the past decade or so, Nvidia has promoted CUDA in schools and research institutes, enabling software such as climate simulation and seismic data processing to be developed based on the platform. Currently, Chinese GPU...

FINDING IV-A: AMD'S AI GPU BUSINESS IS REAL — BUT THE CUDA MOAT IS THE BINDING CONSTRAINT

AMD's Instinct MI300X is not a pretender — it is a genuine computing asset that several of the world's most sophisticated AI operators (Microsoft, Oracle, Meta) have chosen to deploy at scale. The 192GB HBM3 memory architecture provides a structurally differentiated inference advantage for 70B-class models that NVIDIA's H100 cannot match on a per-GPU basis. AMD's Data Center segment revenue of \$12.565B in FY2024 (of which MI300X contributed approximately \$5B+ per Lisa Su's own statements) is real, growing, and profitable. The constraint is not silicon — it is software. NVIDIA's CUDA ecosystem has 15 years of compounding developer lock-in, 4,000+ optimised applications, and the PyTorch default path. AMD's ROCm strategy is the right long-term answer (open-source is the moat-breaker) but the timeline is measured in years, not quarters. The bull case for AMD's AI GPU business is: ROCm achieves enterprise parity by 2026–2027, MI400 hardware competes with Rubin, and AMD takes 15–20% of the AI GPU market by 2028. The bear case is: CUDA lock-in proves sticky, MI400 slips on ROCm readiness, and AMD stabilises at 8–12% market share without a major ecosystem breakthrough.

Evidence: AMD MI300X revenue: \$5B+ annualised run rate (Lisa Su, Q3 2024 earnings) · Azure ND MI300X v5: production since 2024 · OCI MI300X Supercluster: 8,192-GPU cluster deployed · AMD Instinct market share: ~12–13% of AI accelerator revenue (est.) · ROCm GitHub: 1,200+ contributors; 15,000+ stars — but 4× fewer than NVIDIA's CUDA ecosystem by developer engagement metrics

PART V — THE XILINX ACQUISITION: A \$49B BET ON ADAPTIVE COMPUTING

AMD's acquisition of Xilinx, completed February 14, 2022 for \$49B in an all-stock transaction, was the largest deal in AMD's history and one of the three largest semiconductor acquisitions ever executed. The strategic thesis was compelling: Xilinx is the world's leading FPGA (Field-Programmable Gate Array) vendor, with approximately 55% of the global FPGA market and dominant positions in aerospace, defence, automotive, wireless infrastructure, and industrial automation. FPGAs are reconfigurable chips — they can be programmed and reprogrammed for different computational tasks after manufacture, making them ideal for applications where flexibility, low latency, and low power are critical. AMD's CPU+GPU+FPGA+DPU portfolio thesis was that the data centre of the future would be heterogeneous: different workloads requiring different compute substrates, all from a single vendor. Four years into the integration, the Embedded segment's financial performance has been significantly below expectations — but the strategic rationale remains intact.

5.1 Xilinx's Market Position — The FPGA Duopoly

The global FPGA market is a duopoly: Xilinx (now AMD Xilinx) at approximately 55% market share and Intel (Altera) at approximately 30–35%, with a long tail of smaller players (Lattice Semiconductor, Microchip, Achronix). FPGAs serve markets that ASICs cannot serve economically: applications where volumes are too low to justify the \$50M–\$500M NRE (non-recurring engineering) cost of a custom ASIC, or where the design must be updated after deployment. The three largest FPGA markets are aerospace/defence (real-time signal processing, radar, EW systems — AMD Xilinx dominates with security-cleared designs), wireless infrastructure (5G base station digital front-end — AMD Xilinx Zynq RFSoc is the market standard), and data centre acceleration (Microsoft's entire Azure SmartNIC/DPU stack uses AMD Xilinx Alveo devices). The Versal AI Core series (AMD Xilinx's next-gen platform) integrates FPGA fabric with AI Engines — hardened systolic array processors for ML inference acceleration — positioning Xilinx as an AI inference accelerator in addition to a traditional FPGA.

Segment	AMD Xilinx Market Position	Key Products	Est. Revenue (FY2024)	5-Year CAGR Trend
Aerospace & Defence	~60–65% share; top 5 defence primes all use Xilinx FPGAs	Virtex UltraScale+; Versal Premium (space-grade)	~\$1.0–1.2B	+4–6% (stable; defence budget driven)
Wireless Infrastructure (5G)	~55% share; Ericsson, Nokia, Samsung base stations	Zynq UltraScale+ RFSoc; Versal AI Edge	~\$0.8–1.0B	+3–5% (5G rollout dependent)
Data Centre Acceleration	~40% share; Microsoft Azure dominant design win	Alveo U55C; Alveo V80; Versal AI Core	~\$0.6–0.8B	+8–12% (AI inference acceleration)

Industrial & Scientific	~50% share; test equipment, robotics, imaging	Artix/Kintex UltraScale; Versal Prime	~\$0.8–1.0B	+5–7% (IoT + automation)
Automotive (ADAS/Autonomous)	~30% share; competing with Renesas, NXP, Mobileye	Versal AI Edge (automotive grade); Zynq	~\$0.3–0.5B	+12–18% (ADAS acceleration)
Communications (other)	~45% share; telecom, broadcast, network	Virtex/Alveo for network acceleration	~\$0.3–0.4B	+3–5%

SOURCE: Gartner FPGA market share analysis 2024 · Dell'Oro Group wireless infrastructure FPGA market · AMD Embedded segment disclosures 2022–2024 · AMD Investor Day 2022 (Xilinx integration thesis)

5.2 The Embedded Segment Revenue Collapse — What Happened

The Embedded segment (which houses AMD Xilinx FPGA + the semi-custom SoC business from the original AMD Embedded) delivered \$5.8B in FY2023 at its peak — the first full year of Xilinx integration. By FY2024, Embedded revenue had collapsed to \$3.726B: a 36% decline in a single fiscal year. By contrast, AMD's Data Center segment grew 94% in the same period. The collapse was not a market share loss — Xilinx maintained its FPGA leadership positions. It was a double-barrelled inventory digestion event: first, the channel had over-ordered Xilinx FPGAs during the 2021–2022 chip shortage, creating 12–24 months of excess inventory that customers worked through in 2023–2024 without placing new orders; second, the 5G base station deployment cycle in China — Xilinx's largest wireless infrastructure market — slowed materially as Huawei's domestic FPGA alternatives gained market share and Chinese operators prioritised CATV-era equipment. Both headwinds are now largely resolved. AMD guided for Embedded recovery in 2025 as inventory normalises and 5G infrastructure spend resumes in non-China markets.

AMD Embedded Segment Revenue Trajectory:
FY2022 (partial year — Xilinx acq. closed Feb 14, 2022): ~\$3.5B
FY2023 (first full Xilinx year): \$5.766B — Peak
FY2024: \$3.726B — Trough (-35.4% YoY)
FY2025E (AMD guidance): \$4.5–5.5B (recovery; inventory cycle complete)
FY2026E (consensus): \$5.5–6.5B (new 5G cycle + Versal AI adoption)
Root cause of FY2024 trough:
(1) Channel inventory correction: 12–18 months excess burn-through
(2) China 5G slowdown: Huawei FPGA alternatives gaining domestic share
(3) Industrial capex pause: macroeconomic softness in factory automation
Recovery thesis: inventory normalised by Q4 2024; Versal AI Edge design wins
ramping; non-China 5G (India, Southeast Asia, Europe) accelerating.
<i>AMD Embedded guidance FY2025: 'meaningful recovery expected' (Lisa Su, Q4 2024 earnings); specific range not disclosed</i>

5.3 The Xilinx Acquisition Assessment — Was \$49B the Right Price?

The Xilinx acquisition was priced at \$49B in October 2020, when Xilinx was trading at approximately 30× earnings in a semiconductor bull market inflated by COVID-era chip shortages. AMD paid approximately 5.8× Xilinx's FY2021 revenue of \$3.04B — a rich premium even in the

peak of the chip cycle. Four years later, the acquisition is a mixed assessment: the strategic thesis is validated, but the financial execution has underperformed initial expectations due to the inventory digestion cycle and the Embedded trough.

Acquisition Dimension	Assessment	Evidence	Verdict
Strategic thesis (heterogeneous compute)	VALID — data centre increasingly requires CPU+GPU+FPGA+DPU	Microsoft Azure SmartNIC uses Xilinx Alveo; AMD 'adaptive computing' positioning strengthened	✓ CONFIRMED
FPGA market leadership retained	VALID — AMD Xilinx maintains ~55% FPGA market share	No major design win losses; Versal AI Core winning new data centre acceleration contracts	✓ CONFIRMED
Revenue synergies (cross-selling)	PARTIAL — limited cross-sell evidence to date	No publicly disclosed AMD CPU + Xilinx FPGA combined customer wins at scale	~ PARTIAL
Purchase price (\$49B, 5.8× revenue)	EXPENSIVE — paid peak cycle multiple	Xilinx FY2024 revenue (\$3.726B) → implied EV/revenue at acquisition: 13× trough; should recover to 7–8× at \$6B revenue	✗ OVERPAID (cycle-adjusted)
Embedded trough recovery	IN PROGRESS — FY2025 expected recovery	AMD guided for 'meaningful recovery' in FY2025; inventory channel confirmed normalised by Q4 2024	~ WATCH
AI Engine differentiation (Versal)	PROMISING — Versal AI Core winning inference acceleration	Versal AI Core series: hardened AI Engines (400+ TOPS per device); competing with NVIDIA Hopper for specific low-latency inference	✓ EMERGING VALUE
Intangible: talent + IP acquired	HIGH VALUE — 5,000+ Xilinx engineers; 50,000+ IP cores	Xilinx IP library is world's largest FPGA IP repository; no publicly disclosed IP conflicts post-integration	✓ HIGH VALUE
Amortisation impact on GAAP earnings	SIGNIFICANT HEADWIND — \$1.973B annual Xilinx amortisation	FY2024 Xilinx amortisation: ~\$1.97B; depresses GAAP EPS by ~\$1.20/share annually through 2029	✗ FINANCIAL DRAG

SOURCE: AMD Q4 2021 earnings (Xilinx acquisition close announcement) · AMD Embedded segment quarterly disclosures FY2022–FY2024 · AMD Q4 2024 earnings call (Lisa Su Embedded recovery guidance) · Gartner FPGA market share · AMD purchase price allocation (Xilinx \$1.97B annual amortisation from 10-K)

FINDING V-A: XILINX WAS STRATEGICALLY RIGHT AND FINANCIALLY EXPENSIVE — RECOVERY IN PROGRESS

The Xilinx acquisition reflects AMD under Lisa Su at its most ambitious: a cycle-peak valuation paid for a category leader in a market that complements AMD's core CPU+GPU business. The strategic thesis — that data centres need heterogeneous compute including FPGAs — has been validated by Microsoft's Azure Xilinx SmartNIC deployments and the emergence of Versal AI Core as an AI inference accelerator. The financial execution has disappointed relative to \$49B expectations: the Embedded trough (~35% in FY2024) and the \$1.97B annual amortisation charge that suppresses GAAP earnings are real costs. The recovery thesis — inventory cycle complete, Versal AI ramp, non-China 5G acceleration — is credible and already showing in AMD's Q4 2024 sequential Embedded improvement.

By FY2026–FY2027, if Embedded returns to \$6B+ revenue and Versal AI wins data centre inference contracts, the acquisition will appear strategically prescient. If the 5G cycle remains muted and ASIC competition erodes FPGA market share, the \$49B price tag will remain a criticism.

Evidence: Xilinx acquisition cost: \$49B (Oct 2020 announcement; closed Feb 14 2022) · Embedded FY2023 peak: \$5.766B → FY2024 trough: \$3.726B (-35.4%) · Xilinx amortisation: ~\$1.97B annually through ~2028–2029 · FPGA market share: AMD Xilinx ~55%; Intel Altera ~30–35% (Gartner 2024)

PART VI — THE AMD SUPPLY CHAIN: TSMC DEPENDENCY, HBM COMPETITION, AND GLOBALFOUNDRIES

AMD's supply chain is the mirror image of NVIDIA's — almost entirely fabless, almost entirely dependent on TSMC for leading-edge logic, and acutely exposed to the same Taiwan Strait geopolitical risk that defines NVIDIA's existential supply chain vulnerability. Unlike NVIDIA, however, AMD has one additional supplier relationship that neither NVIDIA nor Intel shares: GlobalFoundries, which manufactures AMD's I/O dies and Xilinx FPGA products, providing a partial supply chain hedge that is geographically diversified from Taiwan — though at a cost of less advanced manufacturing for those specific die types.

6.1 TSMC — The 100% Critical Path for AMD's Leading-Edge Products

Every leading-edge AMD product — every Zen 5 CPU compute die, every CDNA3 Instinct GPU die, every Versal AI Core FPGA — is manufactured at TSMC. AMD has no alternative foundry relationship for leading-edge logic. Intel Foundry (IFS) is not a credible alternative for AMD CPU/GPU at AMD's required volumes and yields. Samsung Foundry has not demonstrated TSMC-equivalent yields at N3 or N2. AMD's TSMC dependency is, if anything, more acute than NVIDIA's: AMD uses TSMC across a broader product portfolio (CPUs, GPUs, FPGAs, custom silicon) and is simultaneously one of TSMC's top 5 customers by revenue. A Taiwan Strait disruption scenario that halts TSMC production does not merely halt AMD's AI GPU revenue — it halts AMD's CPU, FPGA, and embedded businesses simultaneously.

AMD Product	TSMC Node	Die Type	Alternative Foundry	TSMC Dependency Risk
Zen 5 CPU (EPYC Turin, Ryzen 9000)	N3 (CCD)	Compute Core Die	NONE at equivalent performance	CRITICAL Ω
CDNA3 GPU (MI300X Instinct)	N5P (GPU tile)	GPU Compute Die	NONE at N5P volume	CRITICAL Ω
Zen 4 CPU I/O Die	GlobalFoundries 6nm	I/O / Interconnect	GlobalFoundries (US + Germany)	MODERATE — GF provides diversification
CDNA3 CPU tiles (MI300X)	N5P	CPU-class die in APU	NONE at N5P	CRITICAL Ω
Versal AI Core (FPGA+AI)	TSMC N7 / N5	FPGA + AI Engine die	NONE at these nodes	HIGH
Xilinx FPGA (legacy Virtex/Kintex)	TSMC 16nm / 7nm	FPGA Fabric	Limited Samsung/Intel for older nodes	MODERATE
Radeon RX GPU (consumer)	TSMC N5 / N4	GPU Compute Die	NONE at volume	HIGH
Custom FPGA (Microsoft Azure Alveo)	TSMC N7	FPGA die	None at N7 volume at AMD's scale	HIGH
AMD Custom APU (PlayStation/Xbox)	TSMC N7 / N6	Custom SoC	TSMC only (locked contracts)	CRITICAL (contractual)

SOURCE: AMD foundry relationship disclosures (AMD 10-K FY2024 risk factors) · GlobalFoundries foundry agreement (AMD/GF relationship disclosed in AMD filings) · TSMC customer list analysis (industry estimates; AMD not officially disclosed by TSMC) · AMD technology roadmap (Zen 5 N3, CDNA3 N5P, CDNA4 N3 confirmed in analyst days)

- GlobalFoundries I/O Die: AMD's I/O dies (memory controllers, PCIe, Infinity Fabric) are made at GlobalFoundries — providing meaningful supply chain diversification from Taiwan for that specific die type
- GlobalFoundries fabs: Fab 8 (Malta, New York, US) + Fab 1 (Dresden, Germany) + Fab 7 (Singapore) — all outside Taiwan; US and EU CHIPS Act beneficiaries
- The I/O die distinction: GlobalFoundries 6nm is NOT competitive with TSMC N3 for compute — AMD uses GF specifically because I/O logic does not need the most advanced node, making GF economics superior for that application
- TSMC Arizona: AMD has confirmed participation in TSMC's Arizona Fab 21 capacity — but Arizona is ramping N4P (not N3) through 2026; AMD's N3 production is entirely Taiwan-based through at least 2027
- Taiwan concentration: approximately 90%+ of AMD's advanced-node production capacity is in Taiwan through 2027, essentially identical to NVIDIA's Taiwan concentration

6.2 HBM Supply — Competing With NVIDIA for the Same Pool

AMD's MI300X requires HBM3 (and MI350X requires HBM3E) — the same High Bandwidth Memory that NVIDIA's H100, H200, and B200 products require. HBM is manufactured by three suppliers globally: SK Hynix (approximately 62% market share), Samsung (approximately 25–30%), and Micron (growing from near-zero in 2023 to approximately 8–12% in 2025). The critical supply chain tension: NVIDIA, as the dominant AI GPU vendor by volume, receives allocation priority from SK Hynix — HBM3E allocated to NVIDIA GB200 production takes precedence over AMD MI350X ramp allocations. AMD competes for the remaining HBM capacity. This creates a structural asymmetry: even if AMD's MI350X hardware is excellent, HBM supply constraints can limit AMD's ability to ramp production in step with demand. Lisa Su acknowledged on the Q3 2024 earnings call that HBM supply constraints had affected MI300X shipment timelines — the same constraint that affected NVIDIA's GB200 ramp.

HBM Supply Competition — AMD vs NVIDIA:

Global HBM market (2025): ~\$30–40B (HBM3/HBM3E)

SK Hynix HBM3E allocation priority: NVIDIA (by volume) → AMD → Others

MI300X HBM3 per unit: 192GB = 6× 32GB HBM3 stacks

GB200 GPU tile HBM3E per unit: 192GB = 6× 32GB HBM3E stacks

GB200 NVL72 rack: 72 GPUs × 192GB = 13,824GB HBM3E per rack

At 100,000 GB200 NVL72 racks produced:

NVIDIA HBM requirement: 1.38 petabytes HBM3E

AMD MI300X at 500,000 units: 96 petabytes HBM3 (lower per-unit count)

SK Hynix capacity constraint: ~300–400M GB HBM3E equivalent annual (2025 est.)

Both NVIDIA and AMD compete for this constrained pool.

NVIDIA's volume advantage = priority allocation = AMD supply disadvantage.

HBM capacity estimates from TrendForce HBM market report Q4 2024; SK Hynix allocation priority order estimated from industry analyst consensus, not publicly disclosed by SK Hynix

6.3 The Full AMD Supply Chain Map

Tier	Supplier	Role	Geography	AMD Concentration Risk
Foundry (Logic)	TSMC (Taiwan)	100% of Zen 5/CDNA3/Versal AI leading-edge dies	Taiwan (Hsinchu, Tainan)	CRITICAL Ω — zero alternative at scale
Foundry (I/O)	GlobalFoundries	AMD CPU I/O dies; legacy FPGA	US (NY), Germany, Singapore	MODERATE — GF unique for AMD I/O; geographically diversified
HBM Memory	SK Hynix (62%), Samsung (25%), Micron (growing)	HBM3/HBM3E for Instinct GPUs	Korea (SK Hynix, Samsung); US (Micron HBM at Boise, Idaho)	HIGH — SK Hynix priority goes to NVIDIA first
GDDR7 Memory	Samsung, SK Hynix, Micron	Radeon RX consumer GPU memory	Korea, US	MODERATE — more suppliers, more competitive
Substrate PCB	Ibiden (Japan), Unimicron (Taiwan)	GPU+CPU package substrates	Japan, Taiwan	HIGH — same suppliers as NVIDIA; concentrated
Advanced Packaging	TSMC CoWoS (AMD MI300X); TSMC SoIC (3D V-Cache)	HBM-on-GPU stacking; cache-on-compute stacking	Taiwan	CRITICAL — TSMC packaging monopoly same as NVIDIA
OSAT (test/assembly)	Amkor (US/Korea), ASE (Taiwan)	Die packaging, test	US, Korea, Taiwan	MODERATE — some geographic diversity
Board/System Integration	Supermicro, Dell, HPE	AMD EPYC + Instinct server systems	US (design); Taiwan/China (assembly)	LOW — multiple integrators
GlobalFoundries (legacy)	GF Fab 8 (Malta, NY)	Older Zen I/O; FPGA legacy; RF chips	US (Malta, New York)	LOW — AMD-GF long-term supply agreement through 2025+

SOURCE: AMD 10-K FY2024 supply chain disclosures · AMD-GlobalFoundries supply agreement (disclosed in AMD SEC filings) · TrendForce HBM market share Q4 2024 · SemiAnalysis advanced packaging supply chain analysis · AMD packaging technology (TSMC CoWoS for MI300X, TSMC SoIC for 3D V-Cache — confirmed in AMD Investor Day 2023)

FINDING VI-A: AMD'S SUPPLY CHAIN IS ESSENTIALLY NVIDIA'S SUPPLY CHAIN — SAME RISKS, SMALLER VOLUME

AMD's supply chain is structurally near-identical to NVIDIA's: both are 100% TSMC-dependent for leading-edge logic, both compete for the same SK Hynix HBM pool, both use TSMC CoWoS advanced packaging, and both are exposed to the same Taiwan Strait geopolitical scenario. AMD's one meaningful differentiation is the GlobalFoundries relationship for I/O dies — approximately 10–15% of AMD's die volume is fabbed in the US (Malta, NY) and Germany (Dresden), providing a partial supply chain hedge that NVIDIA does not have. But for AMD's highest-value products (Zen 5 CCD, CDNA3 GPU, Versal AI), Taiwan is the only option. AMD is additionally disadvantaged relative to NVIDIA in HBM allocation: SK Hynix prioritises its largest customer (NVIDIA) for HBM3E allocation, leaving AMD to compete for the remaining pool. This supply chain asymmetry means that in any HBM supply constraint scenario, AMD's AI GPU ramp is more affected than NVIDIA's.

Evidence: AMD 10-K FY2024: 'TSMC is our sole source for leading-edge logic' (paraphrase of risk factor) · GlobalFoundries manufacturing agreement: disclosed; covers I/O dies + legacy products · HBM market share: SK Hynix 62%; AMD competes with NVIDIA for allocation · AMD MI300X production: 100% at TSMC (N5P) — zero non-TSMC alternative through 2026

PART VII — COMPETITION RISK MATRIX: THREATS ON EVERY FLANK

AMD faces a uniquely complex competitive landscape: it must simultaneously defend its CPU gains against a recovering Intel, challenge NVIDIA's AI GPU monopoly with inferior software ecosystem, defend FPGA market share from ASIC disruption and Intel Altera, and manage the emerging threat from ARM-based server CPUs from Amazon (Graviton), Ampere Computing, and Qualcomm (Oryon) — none of which are x86 and all of which challenge AMD's core CPU franchise from a completely different architectural direction. No other major semiconductor company competes in as many contested markets simultaneously.

7.1 The NVIDIA Threat — CUDA Moat vs AMD's Open Ecosystem

NVIDIA is AMD's primary threat in the Data Center segment. The CUDA software ecosystem is the key battlefield, as analysed in Part IV. Beyond software, NVIDIA's Blackwell architecture and the forthcoming Rubin (2027) maintain hardware leadership in training throughput (GB200 FP8: ~9,000 TFLOPS vs MI300X: ~2,600 TFLOPS). AMD's response must be on multiple fronts: closing the hardware gap with MI400 (N3, 2026), closing the software gap with ROCm 6.x investment, and leveraging price-performance for inference workloads where MI300X's 192GB memory is the differentiator. The risk: NVIDIA's Rubin generation (2027) will likely extend the training performance gap even as AMD closes the inference gap, creating a permanent two-tier market where NVIDIA dominates training and AMD competes primarily for inference.

7.2 The Intel Threat — Resurgence Uncertain but Structural

Intel under CEO Pat Gelsinger (fired December 2024, replaced by interim co-CEOs David Zinsner and Michelle Johnston Holthaus, then Lip-Bu Tan appointed March 2025) has had the most turbulent strategic period in its history. Intel's 18A process node — the bet-the-company foundry technology — has been delayed and its yield rates have not yet reached commercial viability. Intel Gaudi 3 (AI accelerator) has failed to gain significant enterprise adoption against NVIDIA. However, Intel's structural advantages — Intel 18A if it achieves yield parity, the x86 ISA compatibility, and Intel's \$3–5B annual R&D investment in process technology — make Intel a threat that AMD cannot dismiss. The bear case for AMD's CPU business: Intel 18A succeeds, Intel Core Ultra and next-gen Xeon outperform Zen 5 on IPC, and AMD's market share stabilises at 25–28% without further gains.

7.3 The ARM Threat — Architectural Disruption from Below

The ARM server CPU threat is AMD's most underappreciated medium-term risk. Amazon Web Services' Graviton 4 (ARM Neoverse-based, TSMC N3) delivers competitive performance-per-watt against EPYC Turin for cloud-native workloads. Ampere Computing's AmpereOne (192 cores, TSMC N3) directly targets AMD's core-count leadership. Qualcomm's Oryon architecture (used in Snapdragon X Elite laptops) has demonstrated that ARM can achieve near x86 performance at significantly lower power in client computing. As hyperscalers — AMD's fastest-growing CPU customers — increase their ARM deployments (AWS already at 25–30% of new EC2 instances on

Graviton), AMD's EPYC gains share from Intel but faces ceiling pressure from ARM from above. The countervailing factor: enterprise software compatibility with x86 remains powerful for most workloads; ARM is strongest in cloud-native containerised applications, not the broad enterprise software stack.

Competitive Threat	Competitor	Threat Level	AMD Vulnerability	AMD Mitigation	Timeline
AI GPU Training Market	NVIDIA (CUDA moat; B200/Rubin)	CRITICAL	ROCm software ecosystem gap; hardware FP8 gap vs B200	MI400 N3 hardware; ROCm 6.x investment; open-source positioning	FY2025–FY2028 — ongoing
AI GPU Inference Market	NVIDIA (but AMD competitive)	HIGH	CUDA ecosystem still dominant even for inference	MI300X 192GB memory advantage; price-performance; Azure deployment	FY2025–FY2027 — AMD partially winning
x86 Server CPU Market	Intel (Xeon; Clearwater Forest; Intel 18A)	MODERATE	If Intel 18A achieves yield parity, IPC gap may close	Chiplet architecture structural lead; Turin 192 cores; price-performance	FY2026–FY2028 — Intel execution risk high
Cloud Server CPU	ARM (AWS Graviton 4; Ampere AmpereOne; Qualcomm)	HIGH (long-term)	ARM performance-per-watt for cloud-native; hyperscalers building ARM in-house	x86 software compatibility moat; Zen 5 IPC competitive; enterprise sticky	FY2027–FY2030 — structural risk building
FPGA Market	Intel Altera; custom ASIC displacement	MODERATE	Embedded trough; ASIC economics improving for higher-volume FPGA-served apps	Versal AI Engine differentiation; 55% market share; design win backlog	FY2025–FY2028 — recovery expected
Consumer GPU Market	NVIDIA (RTX 5000 series); Intel Arc Battlemage	MODERATE	DLSS 4 vs AMD FSR 4; CUDA gaming ecosystem; NVIDIA brand premium	Radeon RX 9070 XT (strong price-performance); AMD FSR 4 open-source	Ongoing — AMD competitive but not market leader
Custom Silicon (Hyperscalers)	Google TPU v5p; Microsoft Maia; Amazon Trainium	HIGH (long-term)	Hyperscalers building own AI silicon → reduces NVIDIA + AMD TAM	AMD Instinct for scale-out; ROCm compatibility with custom silicon	FY2027–FY2030 — structural TAM risk
Laptop/Client CPU	Intel Core Ultra (Arrow Lake, Lunar Lake); Qualcomm Snapdragon X	MODERATE	Intel recovering in thin-and-light; ARM battery life advantage	Ryzen AI 300 series (TOPS leadership); AMD's NPU roadmap	FY2025–FY2027 — competitive

SOURCE: Mercury Research CPU market share data 2024 · MLCommons AI accelerator benchmark suite 2024 · AWS Graviton 4 launch (Nov 2023); Ampere AmpereOne 192-core (2024) · Intel 18A yield status (Intel Foundry Investor Day 2024; 'world-class yields by 2027' commitment) ·

[NIKKEI] Query: "AMD AMD Instinct MI300X NVIDIA competition GPU market s..." Score:63.3

nd the capacity, about over \$10 billion for the ecosystem [in Taiwan] ... If you think about all of the advanced packaging substrates, test capacity, rack scale design and integration ... it's also very, very important for AMD's long-term growth," Su said. Securing supply chain capacity, she added, is crucial for competing in the next-generation AI race. "I came here to Taiwan to make sure that we prepare ourselves for the significant ramp[-up] in CPUs, and we're getting exce...

FINDING VII-A: AMD IS WINNING THE CPU BATTLE BUT FIGHTING FOR ITS LIFE IN AI GPUS

AMD's competitive matrix reveals a bifurcated story. In server CPUs, AMD is winning — EPYC Turin's architectural superiority (192 cores, 12-channel memory, chiplet economics) gives AMD a structural advantage that is sustainable through at least Zen 6 (2026–2027). Intel's recovery trajectory under new leadership is uncertain and execution-dependent. The ARM threat is real but measured in years, not quarters. In AI GPUs, AMD is a credible challenger but faces a determined adversary with 15 years of software ecosystem investment that cannot be closed through hardware alone. The medium-term outcome depends almost entirely on ROCm's maturation velocity and whether Microsoft/Meta/Oracle's continued MI300X deployment generates sufficient ecosystem momentum to tip the PyTorch ROCm backend from 'supported' to 'preferred' for specific workload classes. The ARM threat to AMD's CPU business is the most underappreciated long-term risk — hyperscalers are the fastest-growing segment of AMD's Data Center CPU revenue, and AWS's Graviton strategy is a structural challenge to AMD's cloud-native CPU share.

Evidence: EPYC Turin server share: ~25–30% Q4 2024 (Mercury Research) · Intel 18A: 'world-class yields by 2027' (Intel Foundry Day 2024) · AWS Graviton 4: ~25–30% of new AWS EC2 instance launches (est.) · AMD MI300X revenue: \$5B+ annualised (Lisa Su Q3 2024) · AMD Radeon vs NVIDIA RTX consumer GPU: NVIDIA 75–80% discrete GPU revenue (Jon Peddie Research Q4 2024)

PART VIII — BUSINESS MODEL VALIDITY: IS AMD'S TRANSFORMATION DURABLE?

AMD's transformation from near-bankruptcy in 2014 to a \$250B+ market-cap semiconductor powerhouse with credible positions in CPUs, AI GPUs, and FPGAs is the most significant corporate turnaround in the history of the semiconductor industry. The central question for investors and strategic counterparties: is this transformation structurally durable, or is it dependent on a specific competitive window (Intel's manufacturing stumbles, NVIDIA's software moat being challenged) that could close as Intel recovers and CUDA lock-in deepens? The adversarial balance sheet below presents the prosecution and defense case for AMD's business model validity.

8.1 The Adversarial Balance Sheet

Factor	BULL CASE (Defence)	BEAR CASE (Prosecution)	Weight
CPU Architecture	Chiplet disaggregation is a permanent structural lead; Intel adopting AMD-style tiles = validation; Zen 5 to Zen 6 roadmap clear	Intel 18A could restore IPC competitiveness by 2027; ARM threat growing in cloud-native segment	Moderate bull edge — architecture lead is structural
AI GPU Hardware	MI300X 192GB memory is genuinely differentiated for inference; MI400 (N3, 2026) targets Blackwell-class performance; price-performance advantage real	NVIDIA B200 outperforms MI300X on training by 3–4×; Rubin (2027) will widen gap; hardware alone does not break CUDA lock-in	Slight bear edge — hardware competitive but not decisive
Software Ecosystem (ROCm)	Open-source ROCm is the right long-term answer; PyTorch ROCm official; Meta/Microsoft contributing; enterprise switching cost falls over time	CUDA moat is 15 years deep; 4,000+ apps; 1M+ developers; switching cost is 6–18 months per enterprise stack; ROCm parity is years away	Bear edge — CUDA moat is AMD's most durable competitive disadvantage
EPYC Server CPU Revenue	\$12.565B Data Center FY2024 (+94% YoY); 25–30% server share; Google/Azure/Oracle all cite EPYC as preferred; structural gains not cyclical	25–30% share ceiling possible — x86 ARM substitution + Intel recovery limits upside; Data Center growth rate will moderate from 94%	Bull edge — CPU momentum is real, even if growth rate normalises
Xilinx Integration	FPGA market leader (55% share); Versal AI Engine opens inference acceleration market; recovery in progress for FY2025	\$49B acquisition at cycle peak; \$1.97B annual amortisation drag on GAAP through 2029; Embedded trough deeper than expected	TIE — strategically right, financially expensive
Supply Chain Risk	GlobalFoundries I/O	90%+ Taiwan	Bear edge — same Ω risk

	diversification is a real hedge vs NVIDIA (pure TSMC); TSMC Arizona participation underway	concentration for leading-edge; same Ω risk as NVIDIA for any Taiwan Strait disruption; HBM allocation disadvantage vs NVIDIA	as NVIDIA but with smaller HBM allocation
Financial Trajectory	Non-GAAP net income \$5.433B FY2024; gross margin 52.7% GAAP; R&D \$5.8B (22%) — investing for long-term moat; strong FCF generation	GAAP net income only \$1.641B due to Xilinx amortisation; debt from acquisition; stock-based compensation dilution; valuation (P/E 50–60× GAAP) prices in perfection	Neutral — strong underlying economics, depressed by Xilinx amortisation
Lisa Su Leadership Premium	Transformation from near-bankruptcy to \$250B+ market cap under single CEO is extraordinary; trust in execution is exceptionally high across customers and capital markets	Key-person risk — Lisa Su is AMD; succession planning not disclosed; any Su departure would reprice AMD meaningfully	Bull edge — Su track record is unique asset, but key-person risk is real
Market Structure	Semiconductor duopoly in x86 CPUs (AMD + Intel) is structurally protected; x86 ISA compatibility gives AMD floor demand even in ARM disruption scenario	NVIDIA at \$3.3T vs AMD at \$250B — AMD is fighting a financially superior opponent in AI GPUs; NVIDIA can outspend AMD on ROCm R&D if it chooses	Bear edge — capital asymmetry vs NVIDIA
Valuation	AMD trades at 50–60× GAAP P/E but 40–50× Non-GAAP — reasonable for a company with 94% Data Center growth and clear AI GPU ramp	GAAP P/E 50–60× prices in near-perfect execution of AI GPU ramp + ROCm maturation + Xilinx recovery simultaneously	Neutral — premium multiple; execution risk is the variable

SOURCE: AMD FY2024 10-K Annual Report (EDGAR) · AMD Q4 2024 earnings call (Lisa Su) · Mercury Research CPU market share · SemiAnalysis AMD competitive analysis · Jon Peddie Research GPU market share Q4 2024 · AMD Non-GAAP bridge (earnings release)

8.2 Five-Year Revenue Scenarios (FY2025–FY2030)

AMD's revenue growth trajectory depends on three independent variables that each have significant uncertainty: the pace of EPYC server CPU share gains, the trajectory of MI300X/MI350X/MI400 AI GPU adoption (directly tied to ROCm maturation), and the speed of Embedded segment recovery. The three scenarios below model these as independent drivers with compounding outcomes.

Year	Bear Case	Base Case	Bull Case	Key Bear Driver	Key Bull Driver
FY2024 (Actual)	\$25.785B	\$25.785B	\$25.785B	—	—
FY2025E	\$26–28B	\$29–32B	\$33–37B	MI350X ramp slower; Embedded recovery delayed	MI350X demand exceeds expectations; Embedded normalises faster
FY2026E	\$28–32B	\$35–42B	\$45–55B	ROCm adoption stalls;	MI400 (N3) captures 15%+ AI

				Intel 18A competitive	GPU market; ROCm achieves parity
FY2027E	\$30–35B	\$42–52B	\$60–75B	ARM cloud CPU share accelerates; NVIDIA Rubin dominates training	EPYC 30%+ server share; MI400 in training clusters at scale
FY2028E	\$32–38B	\$50–62B	\$75–95B	AI CapEx pause reduces GPU demand; hyperscaler custom silicon limits AMD TAM	AMD AI GPU 20%+ revenue share; EPYC 35% server CPU
FY2029E	\$35–40B	\$58–72B	\$90–115B	Chinese competitor equivalents (Huawei Ascend) erode AMD's Asia market	Full-stack AMD (CPU+GPU+FPGA) data centre wins at hyperscaler scale
FY2030E	\$38–45B	\$68–85B	\$105–135B	x86 structural decline; ARM + custom silicon > 40% hyperscaler CPU	AMD competes credibly across CPU, GPU, and adaptive compute at \$100B+ revenue

SOURCE: AMD FY2024 actual (10-K) · Wall Street consensus FY2025–FY2026 (FactSet, Bloomberg) · Blue Lotus Research scenario modelling — NOT investment advice · AMD Data Center segment growth drivers: EPYC trajectory + Instinct AI GPU ramp

AMD Valuation Sanity Check (August 2026):

Market Cap (est.): ~\$250B (at ~\$155/share × ~1.61B diluted shares)

FY2024 Revenue: \$25.785B → Price/Sales: ~9.7×

FY2024 GAAP Net Income: \$1.641B → GAAP P/E: ~152×

FY2024 Non-GAAP Net Income: \$5.433B → Non-GAAP P/E: ~46×

Base Case FY2026E Revenue: \$35–42B

At 8× P/S → Market Cap \$280–336B (modest premium from current)

At 40× Non-GAAP P/E (FY2026E Non-GAAP NI ~\$8–10B) → Market Cap \$320–400B

Bull Case FY2026E Revenue: \$45–55B

At 10× P/S → Market Cap \$450–550B (80–120% upside from current)

Bear Case FY2026E Revenue: \$28–32B

At 7× P/S → Market Cap \$196–224B (slight downside from current)

Verdict: AMD is fairly valued to modestly undervalued in base case; significantly

undervalued vs bull case; marginally overvalued in bear case.

Market cap, share price, and EPS are estimates as of August 2026; AMD does not disclose forward guidance beyond one quarter; all scenarios are illustrative, not financial advice

FINDING VIII-A: AMD'S BUSINESS MODEL IS DURABLE — WITH ONE CRITICAL CAVEAT

AMD's fabless design model, chiplet architecture, and diversified portfolio (CPU+GPU+FPGA) represent a genuinely durable business model that has been stress-tested through near-bankruptcy (2014), competitive irrelevance (2016), and the industry's most significant technology transition (AI GPU) in a decade. The CPU transformation is proven: EPYC Turin's architectural superiority translates to real market share and growing revenue. The Xilinx integration is recovering from an expected post-acquisition inventory cycle. The business model validity question hinges on one critical caveat: AMD's Data Center growth is increasingly dependent on AI GPU adoption that requires ROCm to achieve enterprise software parity. If ROCm achieves first-class PyTorch/vLLM/FlashAttention status by 2026 — which requires sustained engineering investment and enterprise adoption momentum from Microsoft/Oracle/Meta — AMD's AI GPU business will grow to 15–20% market share and the bull case is credible. If ROCm remains second-class, AMD's AI GPU revenue stabilises at \$8–12B annually and NVIDIA consolidates its position. This single software variable — more than any hardware specification — determines whether AMD's \$250B valuation is a bargain or a premium.

Evidence: AMD Data Center FY2024: \$12.565B (+94% YoY) · MI300X revenue: \$5B+ annualised · EPYC server share: 25–30% · AMD Non-GAAP NI: \$5.433B (FY2024); GAAP NI: \$1.641B · Xilinx amortisation: ~\$1.97B annual drag · ROCm maturity: PyTorch official backend operational; performance gap on specialised kernels remains

PART IX — FIVE HYPOTHESIS TESTS: THE AMD DEPOSITION VERDICTS

H1: AMD's chiplet architecture advantage is durable through 2030 and cannot be replicated by Intel in time to recover server CPU market share

EVIDENCE FOR:

- ▶ Zen 5 Turin (192 cores, N3 CCD + GF 6nm I/O) demonstrates chiplet economics at full maturity — Intel's equivalent (Clearwater Forest) won't ship until 2025–2026
- ▶ Intel's adoption of its own tile-based architecture (Clearwater Forest, Sierra Forest) is explicit acknowledgment that AMD's approach was correct — but replication requires 3–5 years of manufacturing learning curve
- ▶ GlobalFoundries supply agreement gives AMD a stable, multi-decade I/O die supplier that Intel cannot easily replicate
- ▶ TSMC N3 CCD yield rates are mature (Apple A17 Pro validated N3 in HVM) — AMD enters N3 with proven process, not a new node risk
- ▶ Core count trajectory: AMD's chiplet approach scales to 256+ cores in Zen 6 era; Intel's monolithic approach faces physics limits

EVIDENCE AGAINST:

- ▶ Intel 18A process (Intel Foundry) could achieve TSMC N3-equivalent performance by 2026–2027 if yield targets are met — reducing AMD's node advantage
- ▶ Qualcomm Oryon and ARM Neoverse (AWS Graviton 4) demonstrate ARM ISA can match x86 IPC in server and laptop markets
- ▶ Intel's AMX matrix extensions (Advanced Matrix Extensions) provide a training/inference acceleration advantage in specific enterprise ML workloads
- ▶ If Intel 18A succeeds at scale, Intel can produce tile-based Xeon on its own advanced process — eliminating AMD's 'TSMC access' advantage

QUANTITATIVE TEST: AMD server CPU share: 0% (2016) → 7% (2019) → 25–30% (2025); Net share gain rate: +3–4 ppt/year (2020–2024). If Intel 18A achieves production yield by 2027, AMD share growth rate likely slows to +1–2 ppt/year. ARM substitution risk in cloud: ~25–30% of new AWS EC2 instances on Graviton (2024) — structural ceiling forming.

VERDICT: CONFIRM — AMD's chiplet architecture advantage is structural and durable through at least 2028. Intel's replication timeline (18A at yield by 2027) and ARM's cloud substitution create a meaningful headwind from 2027 onward, but AMD's 25–30% server share base and design win backlog make further loss below 22–25% unlikely through 2030.

Implication: AMD's EPYC business — approximately 70–80% of the Data Center segment — is on a durable trajectory. Revenue growth from CPU alone (\$8–10B by FY2026E) is sufficient to justify AMD's current valuation. The AI GPU upside is option value on top of a strong CPU foundation.

H2: AMD's MI300X and MI350X will capture a sustainable 15–20% of the AI GPU market by FY2028

EVIDENCE FOR:

- ▶ MI300X 192GB HBM3 provides a structurally differentiated advantage for 70B-class inference — Microsoft, Oracle, Meta have all deployed at scale
- ▶ MI350X (288GB HBM3E, 2025) extends the memory advantage to 140B+ models — directly targets the LLM inference sweet spot

- ▶ AMD's pricing is 20–40% below NVIDIA equivalent for comparable workloads — compelling ROI for price-sensitive enterprise buyers
- ▶ PyTorch ROCm official backend + Meta/Microsoft ROCm engineering contributions reduce the switching cost barrier over time
- ▶ MI400 (N3, 2026) targets NVIDIA Blackwell-class performance — if hardware closes the gap, software gap becomes the only barrier

EVIDENCE AGAINST:

- ▶ NVIDIA CUDA moat: 4,000+ optimised applications, 1M+ developers — switching cost is 6–18 months per enterprise stack
- ▶ NVIDIA Rubin (2027) will likely extend the training performance gap even as MI400 closes the current gap — perpetual moving target
- ▶ Enterprise standardisation on CUDA: most Fortune 500 AI infrastructure teams are CUDA-standardised; switching carries organisational risk
- ▶ SK Hynix HBM allocation priority for NVIDIA constrains AMD's ability to ramp MI350X/MI400 in step with demand
- ▶ AMD Instinct current market share: ~12–13% by revenue (2024) — the 15–20% target requires continued ROCm maturation and no major NVIDIA ecosystem acceleration

QUANTITATIVE TEST: AMD Instinct AI GPU revenue: ~\$5B FY2024 annualised. If AI GPU market reaches \$200B by FY2028, 15% share = \$30B AMD AI GPU revenue. Current trajectory (FY2024 \$5B → FY2025E \$10–12B → FY2026E \$15–20B) is consistent with 15% market share scenario if market grows to \$130–150B. NVIDIA FP8 training advantage: 3–4× over MI300X on H100 comparison — this gap must close for AMD to compete in training (the higher-revenue segment).

VERDICT: PARTIAL CONFIRM — AMD will likely reach 12–15% AI GPU market share by FY2028 in the base case, driven by inference workload wins and continued MI300X/MI350X deployment at tier-1 hyperscalers. The 15–20% target is achievable in the bull case (ROCm parity achieved, MI400 competes on training). Without ROCm enterprise parity, AMD stabilises at 10–13% market share.

Implication: AMD's AI GPU business is not a rounding error — \$5B+ in FY2024 growing at 50%+ CAGR is a material business. But it is a challenger, not an incumbent. The ROCm software investment must be treated as AMD's most critical capital allocation decision through 2026.

H3: The Xilinx \$49B acquisition will generate positive return on investment by FY2028

EVIDENCE FOR:

- ▶ Xilinx FPGA market leadership (55%) is intact — no major design win losses post-acquisition; AMD's distribution and customer relationships have expanded Xilinx's enterprise reach
- ▶ Versal AI Core series (hardened AI Engines + FPGA fabric) opens a new inference acceleration market that Xilinx standalone could not have entered at TSMC's most advanced nodes
- ▶ Embedded segment recovery in progress: inventory cycle complete by Q4 2024; FY2025 guidance indicates meaningful revenue recovery toward \$5B+
- ▶ Microsoft Azure SmartNIC programme (Xilinx Alveo) is a permanent, high-margin design win that validates the data-centre FPGA thesis
- ▶ Intangible value: 5,000+ Xilinx engineers, 50,000+ IP cores, and field-programmable expertise complement AMD's fixed-function GPU/CPU designs

EVIDENCE AGAINST:

- ▶ \$49B purchase price at 5.8× peak-cycle Xilinx revenue was an expensive entry — Xilinx FY2024 revenue (\$3.726B) implies AMD paid 13× trough revenue
- ▶ \$1.97B annual Xilinx amortisation charge suppresses GAAP earnings through approximately 2028–2029 — a real cost of capital
- ▶ Embedded trough was deeper and longer than AMD projected at acquisition close: FY2023 peak \$5.766B → FY2024 trough \$3.726B (-35.4%) is not within the expected variance for a market-leader acquisition
- ▶ ASIC economics are improving for volume FPGA applications — the 'ASIC vs FPGA' threshold is rising, reducing FPGA's addressable market in some high-volume segments
- ▶ ROI calculation: at a 10% cost of capital, \$49B in 2022 dollars requires \$5.9B+ of incremental annual cash flow from Xilinx-attributed revenue to justify the investment; current trajectory requires FY2026–FY2027 revenue recovery to achieve this

QUANTITATIVE TEST: Xilinx revenue at acquisition: \$3.04B FY2021. FY2024 actual: \$3.726B (embedded

segment, which slightly overstates Xilinx contribution). Required revenue by FY2028 for positive ROI (at 10% WACC): approximately \$7–8B Xilinx-attributable revenue. Achievable in bull case (Versal AI acceleration + 5G recovery + industrial rebound); uncertain in base case; not achieved in bear case.

VERDICT: PARTIAL CONFIRM — The Xilinx acquisition is strategically valid and its FPGA market leadership is intact, but the financial ROI case requires a bull case revenue recovery to \$7B+ by FY2028 that is not yet assured. The Embedded trough was a significant execution negative vs expectations. ROI positive by FY2028 is possible (45%) but not the base case (35%).

Implication: The Xilinx acquisition is not an overhang on AMD's strategic thesis — it is part of the thesis. But investors pricing in full Xilinx ROI by FY2026 are too aggressive. The \$1.97B annual GAAP amortisation charge makes GAAP EPS comparisons to Intel misleading; Non-GAAP analysis is essential for fair assessment.

H4: AMD's TSMC foundry dependency represents a systemic risk that is materially underpriced in AMD's equity valuation

EVIDENCE FOR:

- ▶ AMD's entire leading-edge portfolio (Zen 5, CDNA3/4, Versal AI) is 100% TSMC-dependent — zero credible alternative for advanced logic at AMD's required volumes and performance targets
- ▶ Taiwan Strait conflict probability: 15–22% over a 10-year horizon (per RAND, CSIS, Atlantic Council estimates) — not zero, and not priced into AMD's 46× Non-GAAP P/E
- ▶ AMD's TSMC dependency is worse than NVIDIA's in one respect: AMD needs TSMC for both CPU AND GPU AND FPGA simultaneously — a Taiwan disruption destroys all three business lines at once
- ▶ HBM allocation disadvantage compounds the TSMC risk: in a Taiwan disruption scenario, HBM supply (primarily from Korea) would continue, but AMD would have no logic dies to attach it to
- ▶ Arizona Fab 21 is not at the leading node AMD needs (N3 for Zen 5) and will not be until 2027–2028 at earliest — no material diversification before that date

EVIDENCE AGAINST:

- ▶ Taiwan Strait conflict would also destroy NVIDIA, Apple, Qualcomm, MediaTek, and essentially the entire global advanced semiconductor supply chain — the risk is systemic, not AMD-specific
- ▶ US military deterrence of Taiwan Strait conflict is substantial; TSMC's economic value to Taiwan's security is itself a deterrent
- ▶ AMD participates in TSMC Arizona and has a design-in relationship with GlobalFoundries (I/O dies) — more supply chain diversification than NVIDIA has
- ▶ Insurance against a Taiwan scenario is effectively unavailable — it is a civilisational risk that cannot be hedged at the company level

QUANTITATIVE TEST: AMD market cap: ~\$250B. Implied probability of catastrophic TSMC disruption priced in equity: ~2–3% (standard market discount for geopolitical tail risk). Actual RAND/CSIS 10-year Taiwan conflict probability: 15–22%. If 15% probability × AMD revenue destruction of ~85% in disruption scenario = 12.75% expected value discount not in the share price → approximately \$30–35B of AMD market cap is systemically underpriced for geopolitical Ω risk.

VERDICT: CONFIRM — AMD's TSMC dependency is the highest unpriced systemic risk in its equity. The market applies approximately a 2–3% probability discount to Taiwan geopolitical risk when the academic consensus is 15–22% over 10 years. AMD is at least as exposed as NVIDIA to this risk, and arguably more so given its broader TSMC-dependent product portfolio.

Implication: AMD's equity valuation implicitly assumes Taiwan Strait stability through the company's planning horizon. This assumption is not unreasonable — conflict is unlikely in any given year — but the probability is not negligible. Portfolio managers pricing AMD at 46× Non-GAAP P/E should explicit-price the TSMC Ω risk as a tail scenario with non-trivial probability.

H5: AMD will achieve greater than 20% share of the AI accelerator market by FY2030, driven by ROCm open-source ecosystem maturation and CUDA switching momentum

EVIDENCE FOR:

- ▶ Open-source is historically a moat-breaker in software — Linux displaced proprietary Unix; Android displaced proprietary mobile OS; if ROCm achieves true open-source ecosystem momentum, the CUDA moat can be eroded
- ▶ Meta's Llama open-source LLM strategy creates a natural pull toward ROCm: Meta contributes to ROCm, deploys MI300X, and its open models are the primary test case for ROCm compatibility
- ▶ Token economics (DeepSeek-style inference cost collapse) are driving enterprises toward the cheapest capable compute — AMD's price-performance advantage is more compelling as token revenue margins compress
- ▶ Government/sovereign AI initiatives (EU AI Sovereignty, India AI Mission) are explicitly seeking non-NVIDIA alternatives — AMD's ROCm open-source approach is strategically aligned with sovereignty concerns
- ▶ By FY2028, ROCm will have been in serious investment mode for 6+ years — the ecosystem gap will narrow substantially even if it does not close completely

EVIDENCE AGAINST:

- ▶ 20%+ AI GPU market share requires AMD to win meaningful training cluster deployments — where CUDA's advantage is most pronounced and AMD is currently weakest
- ▶ NVIDIA Rubin (2027) and its successor will likely extend the absolute performance gap even as AMD closes the relative gap on training throughput — the target keeps moving
- ▶ AI labs (OpenAI, Anthropic, Mistral) are deeply CUDA-optimised and have essentially zero ROI motivation to invest in ROCm compatibility for their core research stacks
- ▶ The 15-year CUDA developer community is self-reinforcing: new ML engineers learn CUDA because all existing codebases, tutorials, and jobs require CUDA — ecosystem moats strengthen over time, not weaken
- ▶ Custom silicon (Google TPU, Amazon Trainium, Microsoft Maia) may take market share from both NVIDIA and AMD — limiting AMD's TAM expansion

QUANTITATIVE TEST: AMD AI GPU revenue FY2024: ~\$5B. AI accelerator market FY2024: ~\$90–110B. AMD share: ~5%. For 20%+ share by FY2030 at a \$300B AI GPU market: AMD needs \$60B+ AI GPU revenue — implying AMD AI GPU revenue grows 12× in 6 years. This is a bull case, not a base case. Base case: 12–15% share of \$250B market = \$30–37B AMD AI GPU revenue by FY2030.

VERDICT: PARTIAL CONFIRM — AMD reaching 20% AI GPU market share by FY2030 is theoretically possible (open-source ROCm + MI400/MI500 hardware parity + inference market expansion) but requires execution in every variable simultaneously. The base case is 12–15% market share. The 20%+ scenario requires a specific cascade: ROCm enterprise parity by 2026 → enterprise adoption momentum → AI lab validation → broad training cluster adoption → scale economies. Each step is individually uncertain.

Implication: AMD's AI GPU bull case is real, not fabricated — but it requires a longer time horizon than the market may be pricing. The 2030 scenario is more plausible than the 2026 or 2027 scenario. Investors should think of AMD's AI GPU business as a 5–7 year thesis, not a 1–2 year catalyst.

PART X — SUPERFORECASTING: TEN SCENARIOS FOR AMD 2026–2032

The following ten scenarios apply a superforecasting methodology to AMD's most consequential strategic uncertainties. Each scenario includes a calibrated probability, the primary driver if it occurs, and a falsifier — the observation that would cause us to revise the probability downward. Scenarios are not forecasts; they are probability-weighted descriptions of the outcome space.

SCENARIO 1: AMD ACHIEVES 30–35% X86 SERVER CPU MARKET SHARE BY FY2028 — P=55%

Driver: EPYC Turin (Zen 5, 192 cores) + Zen 6 roadmap maintain price-performance superiority; Intel 18A delayed or underperforms; hyperscaler EPYC design win momentum continues

Falsifier: Intel 18A achieves competitive yields at Xeon scale by Q4 2026; or ARM substitution accelerates to >40% of new hyperscaler CPU deployments

SCENARIO 2: AMD'S AI GPU MARKET SHARE STABILISES AT 12–15% BY FY2028 — NOT THE TRANSFORMATIONAL SHARE GAIN THE BULL CASE PROJECTS — P=40%

Driver: ROCm remains 'supported but not preferred' in enterprise stacks; CUDA ecosystem compounding continues; NVIDIA Rubin (2027) extends training performance gap; AMD wins inference but loses training

Falsifier: ROCm achieves measurable enterprise certification milestones; a major AI lab (Anthropic, OpenAI, Mistral) deploys MI400 at training scale; AMD closes FP8 training gap with Rubin

SCENARIO 3: AMD ACHIEVES 20%+ AI GPU MARKET SHARE BY FY2030 — THE TRANSFORMATIONAL BULL CASE — P=25%

Driver: ROCm open-source ecosystem achieves PyTorch first-class parity; MI400/MI500 hardware closes training performance gap; enterprise inference ROI advantage drives adoption at scale; sovereign AI mandates favour open-source alternatives

Falsifier: NVIDIA Rubin architecture is so performant that AMD's price-performance advantage disappears even on inference; or CUDA introduces open-source features that neutralise ROCm's differentiator

SCENARIO 4: AMD'S EMBEDDED SEGMENT RECOVERS TO \$5.5B+ REVENUE IN FY2025 — P=50%

Driver: Inventory digestion complete by Q4 2024 (AMD confirmed); non-China 5G infrastructure resumes ordering; Versal AI Core design wins at hyperscalers begin contributing revenue; industrial automation recovery

Falsifier: China 5G investment remains suppressed through FY2025; Versal AI Core loses to ASIC alternatives in data centre inference; second inventory correction in industrial segment

SCENARIO 5: AMD ACQUIRES OR MERGES WITH A SOFTWARE INFRASTRUCTURE COMPANY TO ACCELERATE ROCm ECOSYSTEM (>\$10B TRANSACTION) — P=35%

Driver: Lisa Su recognises ROCm software gap as existential constraint on AI GPU growth; acquisition of a company with CUDA-compatible middleware, inference runtime, or ML compiler toolchain

Falsifier: AMD decides organic ROCm investment is sufficient; antitrust environment prevents large acquisition; no suitable acquisition target available at acceptable valuation

SCENARIO 6: TAIWAN STRAIT DISRUPTION MATERIALLY AFFECTS AMD PRODUCTION FOR 90+ DAYS — P=17%

Driver: PLA military action; Taiwan political crisis; TSMC factory operations disrupted or workforce evacuated

Falsifier: US military deterrence effective; Taiwan security situation stable; AMD Arizona fab achieves sufficient volume to partially offset (not achievable by 2028)

SCENARIO 7: ARM-BASED CPUs EXCEED 30% OF NEW HYPERSCALER SERVER DEPLOYMENTS BY FY2028 — P=60%

Driver: AWS Graviton 4 (and 5) performance continues to improve; Ampere AmpereOne and Qualcomm cloud CPUs gain certification; software compatibility improves for containerised cloud-native workloads

Falsifier: Enterprise software ISV certification on ARM fails to generalise beyond cloud-native; x86 compatibility requirement persists for broad enterprise stack

SCENARIO 8: AMD LAUNCHES AMD AI ENTERPRISE SOFTWARE PLATFORM (GPU-AS-A-SERVICE + ROCm MANAGED SERVICES) GENERATING \$1B+ ANNUAL RECURRING REVENUE BY FY2028 — P=45%

Driver: AMD recognises that hardware-only model cannot sustain AI GPU business; acquires or builds software layer; follows NVIDIA AI Enterprise model; hyperscalers pay for managed ROCm optimisation

Falsifier: AMD decides software margins do not justify investment; NVIDIA AI Enterprise pricing power too strong to compete with; ROCm organic matures without commercial software layer

SCENARIO 9: INTEL RECOVERS TO 25%+ X86 SERVER CPU MARKET SHARE BY FY2029 (FROM ~70–75% CURRENT) — P=30%

Driver: Intel 18A achieves competitive yields; Clearwater Forest Xeon matches EPYC Turin on core count; Intel pricing competitiveness returns; enterprise Intel loyalty yields to AMD competitive threat reversal

Falsifier: Intel 18A yields remain below competitive threshold; ARM continues to take share from both Intel AND AMD; AMD Zen 6 maintains core count and memory bandwidth superiority

SCENARIO 10: AMD'S TOTAL REVENUE EXCEEDS \$45B IN FY2026 (BASE CASE) — VALIDATING THE AI GPU + CPU DUAL-ENGINE GROWTH THESIS — P=65%

Driver: EPYC Turin wins continuing at hyperscalers; MI300X/MI350X ramp delivers \$10–12B AI GPU revenue in FY2026; Embedded recovery adds \$4.5–5B; Gaming + Client stable

Falsifier: MI350X ramp encounters HBM supply constraints; AMD Data Center growth decelerates below 40% in FY2025; Intel 18A surprise competitive recovery

CONCLUSION — THE AMD DEPOSITION: VERDICT ON THE WORLD'S MOST CREDIBLE CHALLENGER

AMD's due diligence verdict is one of calibrated credibility. The company that Lisa Su inherited in 2014 — burning cash, losing market share on every front, and widely considered a going-concern risk — has been transformed into a genuine multi-market semiconductor powerhouse with structural advantages in CPU design, a credible AI GPU business built on real deployment wins, and an adaptive computing portfolio through Xilinx that no other company replicates. The CPU transformation is complete and structural: EPYC Turin's chiplet architecture leadership is durable through at least 2028, and AMD's 25–30% server CPU market share represents a permanent redistribution of the x86 ecosystem that will not be easily reversed even as Intel recovers under new leadership.

The AI GPU transformation is real but incomplete. The MI300X is not a marketing exercise — it is a deployed, revenue-generating product at Microsoft Azure, Oracle, and Meta at genuine scale. The \$5B+ annualised revenue trajectory in FY2024 growing at 50%+ CAGR is material. But the CUDA software moat remains AMD's most durable competitive disadvantage. NVIDIA's 15-year investment in developer tooling, optimised libraries, and the PyTorch-to-CUDA default path cannot be overcome through hardware specifications alone. The ROCm open-source strategy is the right answer to the right problem — and the open-source approach (MIT/Apache licensing) is structurally superior to CUDA's proprietary model for long-term ecosystem sustainability. But 'right long-term answer' and 'near-term revenue catalyst' are different claims. ROCm's timeline is years, not quarters, and investors pricing in 20%+ AI GPU market share by 2026 are too aggressive. The 15–20% scenario is plausible by 2030. The 12–15% scenario is the more conservative base case by 2028.

AMD's supply chain vulnerability — 90%+ TSMC-dependent for leading-edge logic — is identical to NVIDIA's and equally underpriced in the equity market. The Taiwan Strait risk that applies to NVIDIA applies to AMD with equal or greater force, since AMD depends on TSMC across CPU, GPU, and FPGA simultaneously. The GlobalFoundries relationship for I/O dies provides partial diversification, but it is not sufficient to sustain AMD's business in any disruption scenario that halts TSMC's Taiwan operations. This Ω risk is not unique to AMD — it is a civilisational risk embedded in every advanced semiconductor company — but it is real, non-trivial, and systematically underrepresented in semiconductor equity valuations across the sector.

The AMD deposition's final verdict: AMD is a rare case of a corporate transformation executed with genuine strategic clarity and operational discipline over a decade. Lisa Su's thesis — that AMD could win on architecture, on chiplet economics, on open-source software, and on price-performance for the cost-conscious segment of the market NVIDIA underserves — has been validated in CPUs and is being validated in AI GPUs at a slower pace than bulls hope but at a more credible pace than bears would admit. AMD at ~\$250B market cap is not cheap for a company with \$1.641B GAAP net income and a \$1.97B amortisation drag. But it is not irrational for a company on a trajectory toward \$40–50B revenue by FY2026 with structural positions in the two most contested semiconductor

markets of the decade. The thesis is credible. The execution is proven. The software risk is real. The geopolitical risk is shared. AMD is not NVIDIA — but AMD is, at this moment, the most credible challenger to NVIDIA's AI hardware monopoly that the market has ever seen.

REFERENCE MATRIX — SOURCES AND EVIDENCE

#	Source	Type	Key Data Points Referenced	Section Used
1	AMD FY2024 Annual Report (10-K, EDGAR)	Primary / SEC Filing	Revenue \$25.785B; GAAP GM 52.7%; GAAP NI \$1.641B; Non-GAAP NI \$5.433B; SBC \$1.392B; Xilinx amortisation \$1.97B; segment breakdown	I, II, V, VIII
2	AMD Q4 FY2024 Earnings Call Transcript	Primary / Management	Lisa Su commentary: MI300X \$5B+ run rate; Embedded recovery guidance; Zen 5 EPYC Turin momentum; Data Center +94% YoY	I, II, IV, VII
3	AMD Q3 FY2024 Earnings Call Transcript	Primary / Management	MI300X 'exceeds \$5B revenue run rate'; HBM supply constraints acknowledgment; AI GPU demand commentary	IV, VI
4	AMD Investor Day 2022 (Xilinx Integration)	Primary / Corporate	Xilinx acquisition strategic rationale; FPGA market thesis; Versal AI Engine roadmap; \$49B acquisition price	V
5	AMD Investor Day 2024 (AI GPU Roadmap)	Primary / Corporate	MI300X, MI350X, MI400 roadmap; CDNA4 architecture preview; ROCm investment commitment; data centre growth targets	IV, VIII
6	AMD FY2022–FY2023 Annual Reports (EDGAR)	Primary / SEC Filing	Embedded segment FY2023 peak (\$5.766B); historical segment trajectory; Xilinx purchase price allocation	II, V
7	Mercury Research	Third-Party /	AMD server CPU	III, VII

	CPU Market Share Data Q4 2024	Industry	market share: ~25–30%; Intel Xeon: ~70–75%; historical share trajectory 2016–2024	
8	AnandTech / Chips and Cheese — EPYC Turin Analysis	Technical Media	Zen 5 IPC analysis; chiplet architecture technical breakdown; Turin 192-core benchmark data	III
9	Serve The Home — EPYC Benchmark Database	Technical Media	EPYC Milan/Genoa/Turin server benchmarks vs Xeon; price-performance analysis; memory bandwidth comparisons	III
10	Intel Granite Rapids / Xeon Product Brief (Intel.com)	Primary / Intel	Intel Xeon Granite Rapids specs: 128 cores, 8-channel DDR5, AMX extensions, Intel 3 process	III
11	Intel Foundry Investor Day 2024	Primary / Intel	'World-class yields by 2027'; Intel 18A yield status disclosure; Intel foundry strategy update	III, VII
12	SemiAnalysis — AMD Competitive Analysis (Dylan Patel)	Industry Analysis	MI300X vs H100 performance analysis; CoWoS packaging supply chain; ROCm vs CUDA ecosystem comparison	IV, VI
13	AMD Instinct MI300X Product Brief	Primary / AMD	MI300X specifications: 192GB HBM3, 5.3 TB/s bandwidth, CDNA3 architecture, TSMC N5P process	IV
14	Azure ND MI300X v5 Instance (Microsoft Azure)	Primary / Customer	Microsoft Azure MI300X deployment announcement; production launch 2024; inference workload use case	IV
15	Oracle OCI MI300X Supercluster	Primary / Customer	Oracle MI300X cluster: 8,192 GPUs; training and inference deployment; AMD alternative to NVIDIA validation	IV
16	PyTorch ROCm Documentation (pytorch.org)	Primary / Technical	ROCm official PyTorch backend; supported	IV

			operations; performance gap areas; version parity status	
17	MLCommons AI Inference Benchmark Suite 2024	Industry Benchmark	ROCm vs CUDA performance comparison on standardised inference workloads; MI300X benchmark results	IV
18	Gartner FPGA Market Share Analysis 2024	Industry Research	AMD Xilinx ~55% FPGA market share; Intel Altera ~30–35%; Lattice/others ~15%; market size \$8–10B	V
19	Dell'Oro Group — Wireless Infrastructure FPGA	Industry Research	5G base station FPGA market; Xilinx RFSoc dominant position; China 5G slowdown impact on Embedded	V
20	GlobalFoundries Manufacturing Agreement (AMD 10-K)	Primary / SEC Filing	AMD-GF long-term supply agreement; GF Fab 8 (Malta NY), Fab 1 (Dresden), Fab 7 (Singapore); I/O die manufacturing	VI
21	TrendForce HBM Market Share Report Q4 2024	Industry Research	SK Hynix HBM3E 62% market share; Samsung 25–30%; Micron growing; HBM supply constraint timeline 2024–2025	VI
22	RAND Corporation — Taiwan Strait Analysis (2023)	Academic/Policy	10-year Taiwan Strait conflict probability: 15– 22%; PLA capability assessment; deterrence analysis	IX
23	Jon Peddie Research — GPU Market Share Q4 2024	Industry Research	NVIDIA discrete GPU revenue share: 75–80%; AMD: 12– 17%; Intel Arc: ~3– 5%; consumer + data centre	VII
24	TSMC Technology Roadmap (TSMC Investor Day 2024)	Primary / TSMC	N3/N2/A16/A13 production timeline; CoWoS capacity roadmap; Arizona Fab 21 N4P timeline	VI
25	NVIDIA Blackwell GB200 Product Brief (NVIDIA.com)	Primary / NVIDIA	GB200 specifications; B200 FP8 throughput;	IV

			NVLink interconnect; CoWoS-L packaging	
26	AWS Graviton 4 Technical Documentation (AWS)	Primary / AWS	Graviton 4 TSMC N3; ARM Neoverse V2 core; EC2 instance types; performance vs x86 comparison	VII
27	Ampere Computing AmpereOne 192-Core (2024)	Primary / Ampere	AmpereOne specs: 192 cloud-optimised ARM cores; TSMC N3; power efficiency leadership	VII
28	Lisa Su Harvard Business School Case Study (2023)	Academic / Corporate	AMD transformation narrative 2014–2023; CEO strategy; Bulldozer-to-Zen investment thesis; culture change	I
29	NIKKEI Asia — AMD TSMC Taiwan Supply Chain Coverage	News Media	TSMC-AMD foundry relationship; Taiwan supply chain concentration; N3 production allocation	VI
30	Wall Street Journal — AMD AI GPU Competition Coverage (2024–2026)	News Media	AMD vs NVIDIA AI GPU competition; ROCm vs CUDA enterprise adoption; market share data	IV, VII

SOURCE: All primary-source AMD data from EDGAR public filings (SEC.gov) · Industry data from public sources as cited per row · Blue Lotus Research proprietary intelligence corpus — queries: EDGAR, NIKKEI, WSJ, FT, CNA (BM25 semantic retrieval) — all excerpts represent third-party published information; no material non-public information used.

