

BLUE LOTUS RESEARCH

INTELLIGENCE SERIES -- PhD DUE DILIGENCE

VERTIV: POWERING THE AI AGE

The Critical Infrastructure Landlord of the AI Data Center Supercycle

A PhD-Level Due Diligence Study of Vertiv Holdings (VRT) -- Power & Thermal Infrastructure Business Model, AI Data Center Tailwinds, Competitive Position, and Long-Term Compounding Logic

August 2026 | INTERNAL RESEARCH -- NOT FOR PUBLICATION

Classification: Blue Lotus Research Intelligence Series -- Restricted Circulation

RESEARCH METHODOLOGY

- Qualitative:** Business model anatomy, competitive positioning, management quality assessment.
- Quantitative:** Revenue modeling, margin analysis, data center infrastructure sizing, FCF model.
- Demand:** AI CapEx cycle analysis and power/cooling intensity per GPU generation.
- Superforecasting:** Ten probability-weighted scenarios spanning 2026-2031 AI infrastructure buildout.
- Adversarial:** Bull and bear cases tested against structural assumptions of AI supercycle thesis.

PART I -- VERTIV: THE AI DATA CENTER INFRASTRUCTURE PLATFORM

Vertiv Holdings (NYSE: VRT) is the dominant global supplier of power and thermal management infrastructure for data centers -- the systems that keep servers running, powered, and cool. Unlike the semiconductor companies that capture attention in the AI investment narrative, Vertiv operates at the physical layer of the AI stack: the uninterruptible power supplies (UPS), power distribution units (PDUs), computer room air conditioning (CRAC/CRAH) units, liquid cooling systems, and IT infrastructure management software that every data center -- from hyperscaler AI training cluster to enterprise colocation -- requires to function. Vertiv generated approximately \$8.0B in revenue in FY2024, growing at 17-20% organically, with adjusted EBITDA margins of 18-20% and a backlog of approximately \$7-8B providing exceptional forward revenue visibility. The company is the direct and unavoidable beneficiary of the global AI data center supercycle: every GPU deployed requires power delivery infrastructure, thermal management, and the associated management and monitoring software that Vertiv provides.

The Four Product Pillars

Vertiv organizes its product portfolio into two segments -- Products and Solutions (approximately 75% of revenue) and Services (approximately 25% of revenue). The Products and Solutions segment spans four primary categories that address the full power and thermal management stack of a modern data center.

Product Category	Key Products	Revenue Share (est.)	AI Data Center Growth Driver
Power (AC/DC)	UPS systems, rectifiers, switchgear, busway power distribution	~40% of revenue	AI GPU racks require 30-60kW/rack vs 5-8kW standard -- power density 6-10x
Thermal (Air Cooling)	CRAC/CRAH units, precision cooling, in-row cooling	~25% of revenue	AI clusters require higher airflow density; legacy cooling insufficient for AI racks
Thermal (Liquid Cooling)	Direct liquid cooling (DLC), rear-door heat exchangers, immersion cooling	~12% of revenue	Fastest growing: GB200 NVL72 rack requires liquid cooling; immersion growing

IT Management & Monitoring	DCIM software, power monitoring, remote management	~8% of revenue	AI cluster operations require 24/7 real-time thermal/power monitoring
Services (critical infra)	Preventive maintenance, remote monitoring, field services	~15% of revenue	High-margin, recurring; critical for uptime SLAs in AI training environments

SOURCE: Vertiv FY2024 10-K revenue segment disclosures · Blue Lotus product taxonomy analysis · NVIDIA GB200 infrastructure specs

Why AI Is the Most Powerful Secular Tailwind in Vertiv History

The AI GPU generation transition represents the most significant demand inflection in Vertiv history. Standard enterprise server racks operate at 5-8 kW per rack -- power densities that air cooling handles efficiently with standard CRAC units. An NVIDIA GB200 NVL72 AI training rack operates at approximately 120 kW (up to 140 kW peak) -- 15-25x the power density of a standard server rack. This power density increase has three direct consequences for Vertiv: (1) UPS and power distribution systems must be replaced with higher-capacity units -- existing data center power infrastructure is insufficient for AI racks; (2) Air cooling becomes physically inadequate above 40-50 kW/rack, requiring new liquid cooling infrastructure that Vertiv provides (direct liquid cooling, rear-door heat exchangers, immersion systems); (3) The management complexity of a 120 kW rack requires new monitoring, thermal modeling, and predictive failure analytics that Vertiv provides through its IT management software. Every AI GPU rack installed at a hyperscaler data center generates multiple Vertiv product sales that would not have occurred under traditional server deployment -- and these are incremental to existing infrastructure replacement cycles.

Power Density Comparison: Standard Server vs AI GPU Rack (2026):
Standard 1U server (Xeon/EPYC): 0.5-1.0 kW per server / 5-8 kW per rack
Previous-gen AI server (H100 DGX): 10-15 kW per server / 50-80 kW per rack
Current AI rack (GB200 NVL72): ~120 kW per rack (up to 140 kW peak)
Power density increase (std -> AI): 15-25x increase in kW/rack

Data center power demand growth:
FY2022 total US data center power: ~73 TWh/yr
FY2025 projected: ~120-150 TWh/yr
FY2028 projected (AI scenario): ~200-280 TWh/yr

Vertiv infrastructure required per GB200 NVL72 rack:
UPS capacity: 150-200 kVA module equivalent
Liquid cooling: direct liquid cooling loop + rear-door heat exchanger
Power distribution: high-density PDU + busway segment
Estimated Vertiv infrastructure value per rack: ~\$50,000-120,000

Blue Lotus data center power model · NVIDIA GB200 infrastructure specifications · IEA data center electricity forecast 2025

FINDING I-A: VERTIV IS THE POWER AND COOLING INFRASTRUCTURE TOLL ROAD OF THE AI SUPERCYCLE

Every AI GPU rack deployed globally requires Vertiv infrastructure. At 120 kW per GB200 NVL72 rack -- 15-25x standard server density -- air cooling is physically inadequate, UPS systems must be upgraded, and power distribution requires high-density replacement. Vertiv does not compete at the

silicon layer where NVIDIA and TSMC compete. It competes at the physical infrastructure layer where no AI GPU can be deployed without power delivery and thermal management. Blue Lotus estimates the global AI data center infrastructure upgrade cycle represents a \$50-100B multi-year opportunity for power and cooling infrastructure suppliers, with Vertiv capturing 25-35% as the market leader in critical infrastructure for AI data centers.

Evidence: Vertiv FY2024 10-K · NVIDIA GB200 specifications · IEA data center power forecast · Blue Lotus market sizing

PART II -- FINANCIAL PROFILE: REVENUE, MARGINS, AND BACKLOG FORTRESS

Vertiv has undergone a remarkable financial transformation since its 2020 Platinum Equity SPAC IPO at an enterprise value of approximately \$5B. The company that came public with mid-single-digit organic growth and EBITDA margins below 10% has been completely repositioned: FY2024 revenue of approximately \$8.0B growing at 17-20% organically, adjusted EBITDA margins of 18-20% (expanding rapidly), and a backlog of \$7-8B representing approximately 1.0x annual revenue -- extraordinary visibility by any industrial infrastructure standard. The financial transformation has been driven almost entirely by the AI data center supercycle, which began meaningfully impacting Vertiv orders in FY2022 and has accelerated through FY2023, FY2024, and into FY2025-FY2026.

Revenue Trajectory -- AI Supercycle in Action

Fiscal Year	Revenue	Organic Growth	Adj. EBITDA	Adj. EBITDA Margin	Backlog	Key Event
FY2020	\$4.34B	~2%	~\$390M	~9%	~\$3.0B	SPAC IPO; COVID impact
FY2021	\$4.95B	~7%	~\$480M	~9.7%	~\$3.5B	Recovery; supply chain stress
FY2022	\$5.46B	~10%	~\$580M	~10.6%	~\$4.5B	AI data center orders begin
FY2023	\$6.86B	~22%	~\$950M	~13.8%	~\$6.0B	AI CapEx boom accelerates
FY2024	~\$8.0B	~17%	~\$1.55B	~19.4%	~\$7.5B	Liquid cooling ramp; AI dominates
FY2025E	~\$9.3B	~16%	~\$1.90B	~20.4%	~\$8.5B	GB200 NVL infrastructure ramp
FY2026E	~\$10.8B	~16%	~\$2.35B	~21.8%	~\$9.0B	Continued AI supercycle

SOURCE: Vertiv 10-K FY2020-FY2024 · Blue Lotus FY2025-FY2026E model · Wall Street consensus · Vertiv Q4 FY2024 earnings call

Margin Expansion Story -- The Pricing Power Narrative

The single most important development in the Vertiv financial story beyond revenue growth is the dramatic operating leverage playing out in margins. Adjusted EBITDA margins have expanded from approximately 9% in FY2020 to approximately 19-20% in FY2024 -- a 1,000 basis point improvement in 4 years. This margin expansion has three drivers: (1) price realization -- Vertiv successfully passed through significant price increases in FY2022-FY2023 as supply chain constraints and surging demand gave the company unusual pricing power in a market where customers had no alternative supplier for critical infrastructure with acceptable lead times; (2) operating leverage -- fixed cost infrastructure (engineering, manufacturing footprint, SGA) scales more slowly than revenue as AI demand lifts the top line; (3) mix shift -- liquid cooling and AI-specific power products carry meaningfully higher margins than legacy air cooling for mature data centers. Blue Lotus projects continued margin expansion toward 23-25% adjusted EBITDA by FY2027 as revenue scale and mix shift toward AI-specific, higher-value products continues.

Metric	FY2022	FY2023	FY2024	FY2025E	FY2026E	Driver
Revenue	\$5.46B	\$6.86B	~\$8.0B	~\$9.3B	~\$10.8B	AI infrastructure demand
Adj. EBITDA	\$580M	\$950M	~\$1.55B	~\$1.90B	~\$2.35B	Volume + pricing + mix
Adj. EBITDA %	10.6%	13.8%	~19.4%	~20.4%	~21.8%	+1000bps over 4yr
Adj. EPS (diluted)	~\$0.90	~\$1.42	~\$2.40	~\$3.10	~\$3.90	Operating leverage
FCF Conversion	~50%	~60%	~75%	~80%	~85%	Working capital improvement
Net Debt/EBITDA	~4.5x	~3.0x	~1.8x	~1.2x	~0.8x	Rapid deleveraging
Backlog	\$4.5B	\$6.0B	~\$7.5B	~\$8.5B	~\$9.0B	AI order acceleration

SOURCE: Vertiv FY2022-FY2024 10-K · Blue Lotus FY2025-FY2026E forecast · earnings call guidance

FINDING II-A: VERTIV MARGIN EXPANSION IS STRUCTURAL AND FAR FROM COMPLETE

The transition from 9% to 19% adjusted EBITDA margins in 4 years has been extraordinary but is not yet complete. Vertiv management has guided to adjusted EBITDA margins of 20-22% by FY2025 and has indicated a long-term target of 25%+ as mix shifts toward liquid cooling and AI-specific infrastructure. The path to 25%+ margins involves: (1) continued mix shift -- liquid cooling gross margins are 5-8 percentage points higher than legacy air cooling; (2) services revenue scale -- the \$1.3-1.5B services revenue carries 30-35% adjusted EBITDA margins vs 15-17% for hardware products; (3) procurement scale -- as Vertiv purchase volume scales with AI demand, component cost leverage improves. Blue Lotus projects 23-25% adjusted EBITDA by FY2027-FY2028, implying 2.5-3.0B of EBITDA on 10-12B revenue.

Evidence: Vertiv Q4 FY2024 earnings call guidance · Blue Lotus margin model · liquid cooling ASP analysis

PART III -- LIQUID COOLING: THE HIGHEST-GROWTH PRODUCT IN VERTIV HISTORY

Liquid cooling is the single most important product transition in Vertiv history, and it is being driven entirely by AI GPU density. The physics of heat dissipation are unambiguous: air has a specific heat capacity of 1.006 J/(g*K); water has 4.18 J/(g*K) -- 4.15x more thermal energy absorption per gram. For AI GPU racks operating at 120-140 kW, air simply cannot remove heat fast enough to maintain operating temperatures, regardless of airflow volume. The transition to liquid cooling is therefore not optional for AI data center operators -- it is a physical requirement dictated by GPU power density. Vertiv is the market leader in data center liquid cooling with a product portfolio spanning direct liquid cooling (DLC) with cold plates on GPU/CPU dies, rear-door heat exchangers (RDHx), and full immersion cooling systems where servers are submerged in dielectric fluid.

The Three Liquid Cooling Architectures

Cooling Architecture	How It Works	Power Density Supported	Market Status	Vertiv Product
Air Cooling (legacy)	CRAC/CRAH units circulate chilled air through raised floor or overhead	<50 kW/rack	Incumbent; declining share at leading edge	Liebert DSE/CRV series
Rear-Door Heat Exchanger (RDHx)	Water-cooled door removes heat from server exhaust before room circulation	50-80 kW/rack	Transitional; widely deployed now	Vertiv VRC series
Direct Liquid Cooling (DLC)	Cold plates on GPU/CPU dies with liquid loop; air handles remainder	80-150 kW/rack	Fastest growing; GB200 native design	Vertiv CDU + cold plate systems
Single-Phase Immersion Cooling	Servers submerged in dielectric oil; heat exchanger removes heat	Up to 200 kW/rack	Emerging; used in HPC/AI pilot programs	Vertiv immersion tank + CDU
Two-Phase Immersion Cooling	Dielectric fluid vaporizes on components, condenses in heat exchanger	Up to 400 kW/rack	Early commercial; highest density possible	Partner ecosystem + Vertiv CDU

SOURCE: Blue Lotus cooling architecture taxonomy · NVIDIA GB200 infrastructure guide · Vertiv liquid cooling product portfolio · IDC data center thermal study 2025

The GB200 NVL72 rack -- NVIDIA flagship AI training system -- was designed from the ground up for direct liquid cooling, not air cooling. The rack requires a cooling distribution unit (CDU) delivering chilled water at 18-24 degrees Celsius to cold plates directly attached to GPU dies, with the liquid loop removing approximately 85-90% of heat directly from the chip. Vertiv is a preferred CDU supplier for GB200 NVL72 installations globally. This is not a commodity position -- CDU design must be validated for NVIDIA hardware compatibility, appropriate coolant chemistry, and failure mode safety, creating a certification relationship similar to EDA-foundry certifications. As

hyperscalers deploy GB200 NVL72 racks at scale in FY2025-FY2026, Vertiv CDU revenue is growing at 60-100%+ year-on-year from a small but rapidly scaling base.

Liquid Cooling Market Size and Vertiv Position:	
Global data center cooling market FY2024:	~\$22B total (air + liquid)
Liquid cooling specifically FY2024:	~\$3-4B
Liquid cooling FY2028E (Blue Lotus):	~\$12-18B
Liquid cooling CAGR FY2024-FY2028:	~35-45%
Vertiv liquid cooling revenue (est.):	
FY2023:	~\$400M (6% of Vertiv revenue)
FY2024:	~\$900M (11% of Vertiv revenue)
FY2025E:	~\$2.0B (22% of Vertiv revenue)
FY2026E:	~\$3.5B (32% of Vertiv revenue)
Estimated Vertiv liquid cooling market share: 25-30% globally	
Liquid cooling adj. gross margin: ~35-40% vs ~25-30% air cooling	
Blue Lotus liquid cooling market model · Vertiv investor day disclosures · IDC thermal management forecast 2025	

FINDING III-A: LIQUID COOLING IS THE HIGHEST-MARGIN PRODUCT IN VERTIV HISTORY AND GROWING 60%+ YOY

The transition from air to liquid cooling is an irreversible physics-driven secular shift. Every AI GPU rack above 50 kW/rack -- essentially all current AI training infrastructure -- requires liquid cooling. Vertiv enters this transition as the market leader in data center liquid cooling with validated NVIDIA GB200 compatibility, established CDU manufacturing capacity, and superior gross margins (35-40% liquid vs 25-30% air). As liquid cooling grows from 11% to 30%+ of Vertiv revenue by FY2026, the mix shift alone adds 300-500 basis points to Vertiv's blended adjusted EBITDA margin -- without any organic improvement in underlying product economics. This is one of the most powerful positive gross margin mix shift stories in industrial infrastructure.

Evidence: Vertiv FY2024 investor day · NVIDIA GB200 NVL72 thermal specifications · Blue Lotus liquid cooling model

PART IV -- COMPETITIVE DYNAMICS: VERTIV VS EATON, SCHNEIDER, LEGRAND

The data center power and thermal infrastructure market is competitive but characterized by high switching costs and long qualification cycles that create significant customer retention. Vertiv competes with Eaton (Cooper Industries power products), Schneider Electric (APC brand), ABB, and Legrand across different segments of the infrastructure stack. The competitive landscape varies significantly by product category and customer segment.

Competitor	Primary Strength	Data Center Focus	Est. Market Share	AI Readiness	Threat Level
------------	------------------	-------------------	-------------------	--------------	--------------

Eaton	UPS, switchgear, power quality	30-40% of revenue	~18-22% (UPS)	Moderate -- AI power products launched	HIGH in UPS segment
Schneider Electric (APC)	Full data center stack, EcoStruxure software	~40% of revenue	~22-28% overall	Strong -- EcoStruxure AI integration	HIGH across all categories
ABB	Power distribution, switchgear	<20% data center	~8-12% (power dist.)	Limited AI-specific products	MODERATE
Legrand	PDU, cable management, rack systems	~60% of revenue	~10-15% (PDU/rack)	PDU AI density upgrades launched	LOW-MODERATE
Emerson (spun as Vertiv)	Former parent -- now Vertiv competitor in some areas	N/A (exited)	N/A	N/A	Minimal -- was Vertiv's prior owner
Vertiv (subject)	Critical infrastructure, liquid cooling, AI data center	>75% of revenue	~25-35% (critical infra)	LEADER -- GB200 certified, DLC leader	N/A

SOURCE: Blue Lotus competitive analysis · Vertiv FY2024 10-K · Eaton and Schneider annual reports · market share from IDC Critical Infrastructure market study

Why Schneider Is the Most Formidable Competitor

Schneider Electric is Vertiv primary competitor across the full data center infrastructure stack and has the resources, product breadth, and global service footprint to be a sustained threat. Schneider EcoStruxure platform -- a software layer for data center infrastructure management, energy optimization, and AI-driven predictive maintenance -- is more mature than Vertiv equivalent software offering. Schneider has also moved aggressively into liquid cooling with its EcoAisle containment systems and direct liquid cooling products. The primary structural advantage Vertiv maintains over Schneider is customer concentration and specialization: Vertiv derives 75%+ of revenue from data center customers and has deeper technical relationships with hyperscaler data center architects than Schneider, which must balance industrial automation, building management, and grid infrastructure alongside data center. In AI GPU rack infrastructure specifically, Vertiv engineering teams have validated products for NVIDIA GB200 NVL72 ahead of Schneider, creating a qualification lead in the highest-growth market segment.

Vertiv Competitive Moat -- Service Network and Qualification Lead

Vertiv competitive moat in critical data center infrastructure rests on two pillars: (1) Service network depth: Vertiv operates one of the largest critical infrastructure field service organizations globally, with technicians trained specifically on data center UPS, cooling, and power distribution equipment. Data center operators -- particularly hyperscalers with 99.999% uptime SLAs -- require guaranteed service response times of 4-8 hours globally. Building and maintaining this service network requires years of investment and cannot be rapidly replicated by a new entrant or adjacent competitor. Vertiv service revenue (~\$1.3-1.5B annually) is 30-35% margin and grows proportionally with the installed base -- creating a valuable annuity stream that compounds as more AI infrastructure is deployed. (2) Hardware qualification: hyperscalers qualify data center power and cooling equipment through an extensive technical review process involving thermal modeling, power quality validation, and MTBF/MTTR verification. The qualification process for a critical data center UPS or CDU takes 12-18 months. Once Vertiv is qualified at a hyperscaler, switching requires

repeating this qualification with a competitor -- a process no hyperscaler will undertake unless Vertiv fails on quality or support, creating strong retention dynamics.

PART V -- GEOGRAPHIC AND CUSTOMER CONCENTRATION

Vertiv geographic revenue distribution provides a useful lens on AI data center exposure concentration. The Americas (primarily United States) represent approximately 50-55% of revenue, driven by hyperscaler AI infrastructure investment from AWS, Google, Microsoft, Meta, and Oracle -- the global AI investment leaders. Asia-Pacific represents approximately 20-25% of revenue, with Singapore, Japan, Australia, and increasingly India driving AI data center investment. EMEA represents approximately 25% of revenue, with European hyperscaler and enterprise data center operators as primary customers.

Geography	Est. Revenue Share	Primary Customers	Growth Driver	AI Intensity
Americas	~52%	AWS, Google, Microsoft, Meta, Oracle hyperscalers	AI training cluster infrastructure	HIGHEST -- 70%+ AI-driven growth
EMEA	~25%	European hyperscaler POPs, enterprise colo operators	AI inference + enterprise modernization	MODERATE-HIGH -- growing
Asia-Pacific	~23%	Hyperscaler Asia POPs, Singapore/Japan/India expansion	Sovereign AI data center investment	HIGH and accelerating

SOURCE: Vertiv FY2024 10-K geographic segment · Blue Lotus regional AI infrastructure analysis

Customer Concentration -- The Hyperscaler Dependency

Vertiv largest customers are global hyperscalers -- AWS, Google, Microsoft, Meta, Oracle -- who collectively represent an estimated 40-50% of Vertiv revenue. This concentration is both a strength and a risk: hyperscalers are expanding AI CapEx at 30-60% annual rates through FY2026, generating extraordinary top-line growth for Vertiv, but a simultaneous pause or moderation in hyperscaler AI CapEx would disproportionately impact Vertiv revenue. The AI CapEx cycle dependency is the single most important risk factor for Vertiv -- not competitive, not regulatory, but cyclical. Vertiv management has mitigated this through backlog discipline: the \$7-8B backlog at year-end FY2024, representing approximately 1.0x annual revenue, provides 12-15 months of forward visibility even if new orders were to pause entirely. This backlog buffer insulates Vertiv against an immediate AI CapEx shock far better than most semiconductor hardware suppliers.

FINDING V-A: VERTIV BACKLOG FORTRESS PROVIDES 12-15 MONTHS OF DOWNSIDE PROTECTION

A \$7-8B backlog equivalent to approximately 1.0x annual revenue is extraordinary for an industrial infrastructure company. For context, industrial companies with best-in-class backlog visibility typically carry 0.4-0.6x revenue in backlog. Vertiv 1.0x+ backlog has been built during the AI CapEx supercycle as hyperscalers pre-committed infrastructure orders with long lead times (12-24 months for UPS systems; 18-30 months for large liquid cooling installations). Even if AI CapEx were to pause completely in H1 FY2025, Vertiv revenue would continue growing through FY2025 and into H1 FY2026 from backlog conversion. This backlog fortress is a crucial buffer that makes Vertiv one of the

most defensive plays in the AI infrastructure theme.

Evidence: Vertiv Q4 FY2024 earnings call · Vertiv 10-K backlog disclosure · Blue Lotus backlog conversion model

PART VI -- RISK MATRIX: THE FIVE THREATS TO THE VERTIV THESIS

Risk Factor	Category	P(Materialise, 3yr)	Revenue Impact	Mitigation
AI CapEx cycle pause (hyperscaler moderation)	Demand/Cyclical	30%	-15-25% growth rate compression	Backlog fortress (12-15mo visibility); services ARR
Liquid cooling competition intensifies	Competitive	40%	-3-5% market share over 3yr	Qualification lead; service network scale
Component supply chain disruption	Operational	25%	-5-10% revenue from delivery delays	Dual-sourcing; strategic inventory buffer post-2022
Macro slowdown / enterprise capex freeze	Macro	25%	Enterprise segment -15-25%; hyperscaler stable	Hyperscaler mix growing; enterprise <25% of revenue
Schneider Electric liquid cooling breakthrough	Competitive	35%	-3-8% in liquid cooling over 3yr	Early qualification; service depth; NVIDIA partnership
Geopolitical disruption (manufacturing)	Geopolitical	15%	-5-10% from supply disruption	Manufacturing footprint in Americas + Europe
AI efficiency reduces power density growth	Technical	20%	Power density growth slows to 2-3x vs 15-25x	Efficiency reduces per-rack spend, not absolute demand
Debt load / refinancing risk	Financial	10%	Interest burden; limits capital allocation	Net debt rapidly declining; FY2026E <1x EBITDA

SOURCE: Blue Lotus VRT risk assessment · Vertiv FY2024 10-K risk factors · industry competitive analysis

The AI CapEx Cyclicity Risk -- The Most Important Bear Case

The primary bear case for Vertiv rests on AI CapEx cyclicity: if hyperscalers moderate or pause AI infrastructure investment in response to a recession, interest rate shock, or AI ROI disappointment, Vertiv order intake would compress dramatically. However, several structural factors mitigate the cyclicity risk more than the headline thesis implies: (1) Backlog insulation: Vertiv \$7-8B backlog provides 12-15 months of revenue visibility even with zero new orders -- giving management time to adjust capacity; (2) Services revenue stability: the \$1.3-1.5B services revenue carries minimal cyclicity -- preventive maintenance on installed data center infrastructure is non-discretionary for operators with uptime SLAs; (3) The AI infrastructure buildout is multi-year and committed:

hyperscaler data center construction (civil, electrical, HVAC) has 24-36 month cycles -- once a data center build is committed, the Vertiv equipment must be ordered regardless of market conditions; (4) Enterprise and colocation: approximately 30-40% of Vertiv revenue comes from non-hyperscaler customers (enterprise, colocation, edge) with different capex cycles that provide diversification.

PART VII -- BUSINESS MODEL VALIDITY: ADVERSARIAL ASSESSMENT

Bull Case -- Vertiv at Maximum Thesis Strength

The strongest bull case for Vertiv rests on the irreversibility of the AI power density problem. At 120 kW per GB200 NVL72 rack, liquid cooling is not a preference -- it is a physical requirement. As NVIDIA Rubin architecture (projected FY2027) and subsequent generations push rack power toward 200-250 kW, the liquid cooling requirement deepens. Vertiv is the market leader in data center liquid cooling with validated NVIDIA compatibility and growing manufacturing scale. Services revenue (\$1.3-1.5B at 30-35% margins) creates an annuity that grows with the installed base regardless of CapEx cycle timing. Margin expansion from 19% to 23-25% adjusted EBITDA by FY2027-FY2028 creates compounding earnings growth that outpaces revenue growth. Blue Lotus bull case FY2028: \$13-14B revenue, 24-25% EBITDA, \$3.1-3.5B EBITDA, ~\$5.50-6.50 EPS at 30-35% FCF conversion. At 30-35x forward EPS: fair value range \$165-228 per share.

Bear Case -- The Strongest Challenges to the Vertiv Thesis

The bear case focuses on three legitimate risks: (1) valuation -- at 35-50x forward earnings in FY2024-FY2025, VRT has traded at significant premiums to industrial infrastructure peers (typically 15-25x), embedding AI supercycle expectations that are vulnerable to any demand disappointment; (2) competitive margin compression -- as Schneider, Eaton, and Asian manufacturers invest aggressively in liquid cooling, pricing pressure could limit gross margin expansion in liquid cooling from the current 35-40% toward 25-30% by FY2027-FY2028; (3) AI efficiency risk -- if AI architectures evolve toward significantly lower power consumption per FLOP (through techniques analogous to DeepSeek algorithmic efficiency), the rack power density growth trajectory could slow, moderating the urgency and scale of liquid cooling investment. Blue Lotus bear case FY2028: \$9-10B revenue, 18-19% EBITDA, \$1.6-1.9B EBITDA, ~\$2.80-3.20 EPS. At 20-25x forward EPS: fair value \$56-80 per share.

PART VIII -- HYPOTHESIS TESTS: FIVE CRITICAL INVESTMENT QUESTIONS

H1: AI GPU power density growth will sustain liquid cooling demand at 35%+ CAGR through FY2028

EVIDENCE FOR:

- ▶ GB200 NVL72 at 120 kW/rack has already crossed the air cooling threshold -- liquid mandatory
- ▶ NVIDIA Rubin architecture (FY2027E) projected at 150-200 kW/rack -- further density increase
- ▶ Global hyperscaler AI CapEx guidance FY2025-FY2026: \$300B+ combined -- infrastructure follows
- ▶ Vertiv liquid cooling revenue growing at 100%+ in FY2024 from validated installed base

EVIDENCE AGAINST:

- ▶ AI algorithmic efficiency (DeepSeek-style) could reduce compute per task, moderating new rack deployments
- ▶ Schneider/Eaton competitive pressure could suppress ASPs, reducing revenue growth below demand growth
- ▶ Manufacturing capacity constraints at Vertiv could limit supply even with demand intact

QUANTITATIVE TEST: Liquid cooling market from \$3-4B (FY2024) to \$12-18B (FY2028) = 35-45% CAGR. Vertiv capturing 25-30% share implies \$3.0-5.4B liquid cooling revenue by FY2028 vs ~\$0.9B FY2024.

VERDICT: CONFIRM -- Physics mandates liquid cooling above 50 kW/rack; AI GPU trajectory confirms continued density increase; demand structural through FY2028.

Implication: Liquid cooling mix shift is the primary margin expansion driver for Vertiv through FY2028. Overweight liquid cooling exposure in Vertiv valuation framework.

H2: Vertiv adjusted EBITDA margins will reach 23-25% by FY2027 on current trajectory

EVIDENCE FOR:

- ▶ FY2020 to FY2024 expansion: 9% to 19% adjusted EBITDA -- 1,000bps in 4 years
- ▶ Liquid cooling mix shift adds 300-500bps (liquid gross margins 35-40% vs air 25-30%)
- ▶ Services revenue scale: \$1.3-1.5B at 30-35% margins; growing with installed base
- ▶ Operating leverage: fixed cost infrastructure scales slowly vs 16-20% revenue CAGR

EVIDENCE AGAINST:

- ▶ Competitive pricing pressure in liquid cooling from Schneider, Asian manufacturers
- ▶ Labor and component cost inflation could compress gross margins
- ▶ Investment in manufacturing scale-up could temporarily suppress margins

QUANTITATIVE TEST: Blue Lotus FY2027E: \$11.5B revenue, 23% EBITDA = \$2.65B EBITDA. Even at 21% (conservative), FY2027E EBITDA = \$2.4B vs ~\$1.55B FY2024 -- 60%+ increase.

VERDICT: CONFIRM -- Margin expansion trajectory is structural and driven by product mix shift, not one-time pricing. 23-25% adjusted EBITDA by FY2027 is achievable with base-case AI demand.

Implication: EPS growth will outpace revenue growth through FY2027 via combined revenue scale and margin expansion. Operating leverage significantly underappreciated by the market.

H3: Hyperscaler AI CapEx will pause or contract in FY2026, creating a Vertiv order shock

EVIDENCE FOR:

- ▶ AI ROI has not been demonstrated at sufficient scale to justify indefinite hyperscaler spend growth
- ▶ DeepSeek-style efficiency improvements reduce compute requirements per task, potentially moderating GPU demand
- ▶ Interest rate sensitivity: high rates increase hyperscaler capital cost, compressing NPV of data center investment
- ▶ Historical precedent: 2001 telecom infrastructure overbuild followed by 60%+ demand collapse

EVIDENCE AGAINST:

- ▶ Hyperscalers have publicly committed FY2025 AI CapEx of \$300B+ -- unprecedented forward guidance
- ▶ AI applications are still in early deployment phase -- corporate adoption wave beginning, not ending
- ▶ Vertiv \$7-8B backlog converts to revenue regardless of order intake in next 12-15 months
- ▶ Sovereign AI (UAE, Saudi Arabia, India, France) provides new CapEx stream beyond Big 4 hyperscalers

QUANTITATIVE TEST: Even a 30% reduction in new AI infrastructure orders in FY2026 would be absorbed by Vertiv backlog through H1 FY2027. Revenue would still grow in FY2026. Revenue contraction is a FY2027 risk, not FY2026.

VERDICT: REJECT for FY2026; PARTIAL CONFIRM as FY2027 tail risk (P=30%). Backlog fortress insulates Vertiv for at least 4-6 additional quarters regardless of order intake.

Implication: VRT remains investable through FY2026 even under an AI CapEx moderation scenario. Position sizing should respect FY2027 order cycle risk.

H4: Vertiv service revenue will reach \$2.5B+ by FY2028, creating a recurring revenue floor

EVIDENCE FOR:

- ▶ Services grows proportionally to installed base -- every Vertiv hardware deployment adds to future service revenue
- ▶ FY2024 services ~\$1.3-1.5B at 30-35% margins; AI data center surge adds 50%+ to installed base in 2 years
- ▶ Service contracts are multi-year with high renewal rates -- uptime SLA requirements make service non-discretionary
- ▶ Premium monitoring and remote management services (software layer) growing at 25%+ with AI infrastructure

EVIDENCE AGAINST:

- ▶ Competition from third-party maintainers (TPM) could erode service pricing
- ▶ Hyperscalers increasingly self-maintain critical infrastructure, reducing Vertiv service attach rates

QUANTITATIVE TEST: FY2024 installed base generating ~\$1.4B services. FY2024-FY2028 hardware revenue adds \$15-20B to installed base. At 8-10% service attach rate: \$1.2-2.0B incremental services by FY2028.

VERDICT: CONFIRM -- Services will reach \$2.0-2.5B by FY2028, creating a \$600-900M annuity floor that limits Vertiv downside in any hardware cycle moderation scenario.

Implication: Services floor creates a significant valuation support for VRT that pure hardware cycle analysis underweights. Floor EPS from services alone by FY2028: ~\$1.80-2.20.

H5: Vertiv is significantly undervalued relative to its structural position as AI infrastructure landlord

EVIDENCE FOR:

- ▶ VRT is the most leveraged public equity to AI GPU rack deployment globally -- zero exceptions
- ▶ Earnings growth trajectory (40-50% EPS CAGR FY2024-FY2026E) exceeds current forward multiple expansion
- ▶ Services annuity and backlog fortress create a durability premium not reflected in industrial infrastructure multiples
- ▶ International expansion (APAC sovereign AI) provides TAM expansion not in consensus models

EVIDENCE AGAINST:

- ▶ At 30-45x forward earnings, VRT is priced as a software company with industrial infrastructure risks
- ▶ Cyclicity of hardware demand creates re-rating risk if AI CapEx moderates
- ▶ Competitive intensity in liquid cooling will pressure gross margins by FY2027

QUANTITATIVE TEST: Blue Lotus FY2026E EPS: ~\$3.90. At 30x: fair value \$117. At 35x: \$136. At 40x: \$156. Current consensus implies 25-30x on FY2026E earnings -- full but not extreme.

VERDICT: PARTIAL CONFIRM -- VRT is fairly valued to modestly undervalued on base case AI CapEx

continuation. Premium multiple is justified by services annuity and backlog durability, not justified against bear case AI moderation.

Implication: VRT is a core holding in AI infrastructure theme with position sizing discipline. Trim on significant re-rating to 45x+ forward earnings; add on weakness toward 20-25x during market corrections.

PART IX -- SUPERFORECASTING: TEN SCENARIOS 2026-2031

SCENARIO 1: AI DATA CENTER SUPERCYCLE CONTINUES -- VERTIV 15%+ CAGR THROUGH FY2028 — P=50%

Driver: Hyperscaler AI CapEx sustains \$300B+ annually through FY2027; GB200/Rubin racks deploy at scale; liquid cooling penetration accelerates; Vertiv captures 25-30% of growing liquid cooling market.

Falsifier: AI CapEx pause; AI efficiency breakthrough reduces compute demand per task; competitive margin compression from Schneider/Asian manufacturers.

SCENARIO 2: LIQUID COOLING BECOMES 35%+ OF VERTIV REVENUE BY FY2027 — P=55%

Driver: Rubin architecture (FY2027) at 150-200 kW/rack makes air cooling impossible across entire AI tier; DLC and immersion adoption rates exceed Blue Lotus base case; ASPs hold due to Vertiv qualification lead.

Falsifier: NVIDIA Rubin architecture achieves power efficiency gains that limit rack density growth; Schneider liquid cooling qualification catches Vertiv; immersion cooling adoption slower than expected.

SCENARIO 3: AI CAPEX MODERATION -- VERTIV ORDER INTAKE -30% FROM PEAK — P=30%

Driver: Interest rate shock; AI ROI disappointment at enterprise scale; hyperscaler earnings pressure forces CapEx rationalization; sovereign AI programs delayed.

Falsifier: AI enterprise adoption accelerates; new AI applications expand TAM; sovereign AI investment backfills any hyperscaler moderation.

SCENARIO 4: VERTIV ADJUSTED EBITDA REACHES 25%+ BY FY2027 — P=40%

Driver: Liquid cooling mix exceeds 30% of revenue; services annuity scales to \$2.0B+; operating leverage on fixed cost base; AI-specific products (high-density UPS, CDU) carry 40%+ gross margins.

Falsifier: Competitive pricing pressure; component cost inflation; investment in manufacturing scale suppresses near-term margins.

SCENARIO 5: VERTIV MAKES A STRATEGIC ACQUISITION IN LIQUID COOLING OR IT MANAGEMENT —

P=35%

Driver: Post-deleveraging (FY2026E net debt <1x EBITDA), Vertiv has capacity for \$2-4B acquisition; targets in IT infrastructure management software, liquid cooling niche players, or DCIM platforms.

Falsifier: Management commits to pure organic growth and buybacks; acquisition targets are overpriced; integration risk deters board.

SCENARIO 6: SCHNEIDER ELECTRIC INTENSIFIES LIQUID COOLING COMPETITION -- VERTIV MARKET SHARE PRESSURE — P=45%

Driver: Schneider \$1B+ liquid cooling investment materializes in qualified products by FY2026; hyperscalers dual-source CDU and DLC systems from both Schneider and Vertiv; ASP compression begins.

Falsifier: Schneider qualification delayed beyond FY2027; Vertiv service network advantage sustains pricing; NVIDIA ecosystem preference maintains Vertiv position.

SCENARIO 7: DATA CENTER POWER BECOMES A BOTTLENECK -- UTILITY GRID DELAYS HURT VERTIV REVENUE TIMING — P=40%

Driver: US utility grid cannot provide power connections to new data centers fast enough; hyperscalers delay equipment orders until power is secured; Vertiv backlog stretches but revenue timing slips.

Falsifier: Nuclear power agreements (Three Mile Island, etc.) accelerate grid capacity; hyperscalers invest in on-site generation; grid modernization proceeds faster than expected.

SCENARIO 8: SOVEREIGN AI DATA CENTER INVESTMENT BECOMES 15%+ OF VERTIV REVENUE BY FY2028 — P=35%

Driver: UAE, Saudi Arabia, India, France, Japan sovereign AI programs commit to national AI compute infrastructure; each program requires \$1-5B in power and cooling infrastructure; Vertiv wins significant share.

Falsifier: Sovereign AI programs delayed by political risk; Chinese infrastructure suppliers win Asia/Middle East share; US export controls restrict Vertiv to limited markets.

SCENARIO 9: VRT STOCK DOUBLES FROM FY2024 LEVELS BY FY2027 (TARGET: \$120-150+ RANGE) — P=35%

Driver: EPS grows 40-50% CAGR to \$5.50-6.50 by FY2027E; multiple re-rates from 25x to 30-35x on demonstrated durable AI infrastructure thesis; buybacks accelerate post-deleveraging.

Falsifier: AI CapEx moderation compresses earnings growth; multiple contracts to 15-20x on cyclical concerns; competition pressure limits margin expansion.

SCENARIO 10: EDGE AI DATA CENTER CREATES NEW VERTIV MARKET (TELCO EDGE, RETAIL EDGE, INDUSTRIAL EDGE) — P=30%

Driver: AI inference at the network edge (telco 5G towers, retail stores, factories) requires compact, reliable power and cooling infrastructure; Vertiv modular micro-data center products address this market.

Falsifier: Edge AI inference remains cloud-bound; economics of edge infrastructure unfavorable; market addressed by smaller specialist vendors.

CONCLUSION -- VERTIV: THE UNAVOIDABLE AI INFRASTRUCTURE PLAY

The Vertiv due diligence concludes with a clear structural finding: the company occupies the most defensible position in the AI infrastructure value chain short of the GPU manufacturers themselves. Every AI GPU deployed requires Vertiv power delivery and thermal management -- and as GPU power density has crossed the 120 kW/rack threshold with GB200 NVL72, this requirement has become technically mandatory for liquid cooling rather than optionally preferred.

The financial transformation from 9% to 19%+ adjusted EBITDA margins in four years is not complete -- Blue Lotus projects continued expansion toward 23-25% by FY2027 as liquid cooling mix grows, services scale, and operating leverage compounds. The \$7-8B backlog fortress provides 12-15 months of revenue visibility that shields Vertiv from any near-term AI CapEx cycle moderation, giving management a genuine runway to execute without near-term revenue risk.

The primary risks -- AI CapEx cyclical, competitive intensity in liquid cooling, and valuation premium -- are real and should be respected. Vertiv is not a risk-free compounder; it is a high-quality cyclical with structural AI tailwinds and a services annuity that limits downside. For long-term investors, Vertiv represents one of the most compelling infrastructure plays in the AI supercycle: the company that makes the AI data center physically possible.

REFERENCE MATRIX

#	Source	Type	Key Data Points
1	Vertiv 10-K FY2024 (EDGAR)	Primary / SEC filing	Revenue ~\$8.0B, adj. EBITDA margin ~19.4%, backlog ~\$7.5B, geographic segments
2	Vertiv 10-K FY2020-FY2023 (EDGAR)	Primary / SEC filings	Historical revenue, margin expansion trajectory, product segment data
3	Vertiv Q4 FY2024 Earnings Call	Company disclosure	FY2025E guidance, liquid cooling growth commentary, backlog composition
4	Blue Lotus corpus, EDGAR collection, 139,756 chunks	Internal research corpus	Vertiv financial data, customer concentration, risk factors
5	Vertiv Investor Day 2024	Company disclosure	Long-term margin targets, liquid cooling strategy, services growth

			plan
6	NVIDIA GB200 NVL72 Infrastructure Deployment Guide	Technical specification	120 kW rack power density, CDU specifications, cooling requirements
7	IEA Data Center Electricity 2025 Report	Government/research	Global data center power consumption growth, AI-driven demand projection
8	Blue Lotus corpus, WSJ collection, 335,496 chunks	Internal research corpus	AI data center infrastructure investment, competitive dynamics
9	Blue Lotus corpus, NIKKEI collection, 397,454 chunks	Internal research corpus	APAC data center build-out, hyperscaler investment, power demand data
10	IDC Critical Infrastructure Market Study 2025	Industry research	Market share data, competitive landscape, liquid cooling market sizing
11	Schneider Electric Annual Report 2024	Company disclosure	EcoStruxure platform, data center segment competitive position
12	Eaton Power Quality Division Annual Report 2024	Company disclosure	UPS market share, AI data center product strategy
13	Blue Lotus corpus, FT collection, 109,507 chunks	Internal research corpus	AI infrastructure boom, data center investment, power grid constraints
14	US Department of Energy AI Data Center Report 2025	Government	US data center power demand growth, grid capacity constraints
15	Blue Lotus Vertiv margin and scenario model FY2024-FY2028	Internal analysis	Revenue and EBITDA scenarios, liquid cooling sizing, competitive assessment